



KTV Tree: Interactive Top-K Aggregation on Large Dataset in Cloud


Yuzhe Tang, Ling Liu, Junichi Tatemura, Hakan Hacigumus



Introduction on Top-K Aggregation

- Query: Top-K aggregation with range selections
- App: What are tweeted now in a particular area?
 - Schema: Recent Tweets (tag, *long*, *lat*, ...)
 - Top-k popular tags in a given geographic area.

```
SELECT tag, count(*) as c FROM tweets
WHERE long BETWEEN I1 AND I2
      lat BETWEEN I3 AND I4
GROUP BY tag
ORDER BY c LIMIT k
```






Problem Statement


- For schema $\langle a, b \rangle$, given queried range $[b1, b2]$, find top- k popular records within limited latency.

Aggregated popularity



1D range \leftarrow 2D range

Hilbert space-filling curve



Related Work (on Top-K)

- Top-k query processing algorithms
 - FA (pods96), TA (pods01, jccs03), TPUT (podc04), KLEE (VLDB05)
- Materialized View for Top-k Query
 - Yi Ke (icde03), LPTA (vldb06), Top-k monitoring (sigmod03)
- Few prior work addresses “**interactive** top-k aggregation processing with **dynamic range predicates**.”




Baselines

- Partitioning: Given schema $\langle a, b \rangle$, we range-partition data on attribute b .
- Two baselines for top-k processing
 - Local view with threshold algorithm (TA)
 - Segment tree-based view



Baseline 1: Local View




Problem of Local Views

- Can't scale to large # of partitions in query latency and costs.

Baseline 2: Segment-tree View

a	c
g1	28
g3	16
...	...



Problem of Tree Views

- *Extra maintenance overhead.*
- *High-level views handle global updates, leading to bottleneck.*

Benefits of Segment Tree Views


- *For query spanning r leaves, only $\log(r)$ internal nodes are required for query answering*

KTV-Tree: Threshold-based Incomplete View Tree



- Basic idea:
 - Threshold on each node to filter out updates on small values.


- The maintained view is incomplete, due to the threshold filter out certain updates.






Incremental tree maintenance

- Update the views (given fixed threshold).
 - Given updates from child, decide whether to report to parents.
 - Based on threshold
- Update the thresholds.
 - Triggered periodically



Update the thresholds

- Step 1: a top-down process for updating threshold
 - Initiated by root node,
 - Propagate down to leaf, such that $T_{\text{parent}} = \sum T_{\text{child}}$.
- Step 2: a bottom-up process to refill view entries.



Experiment setups

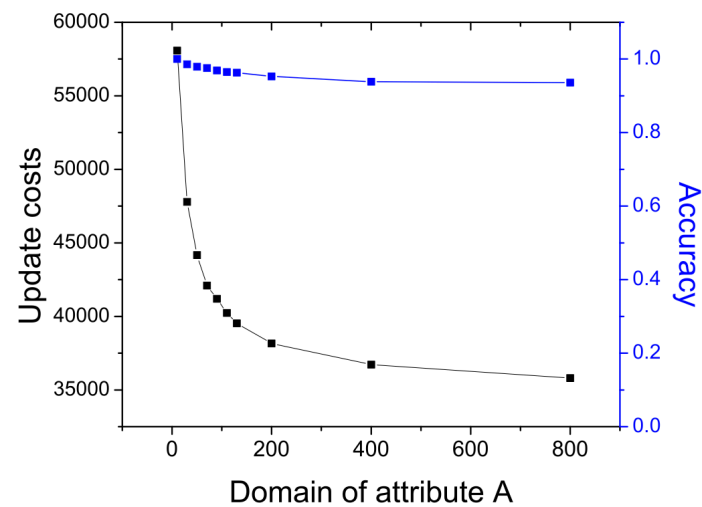
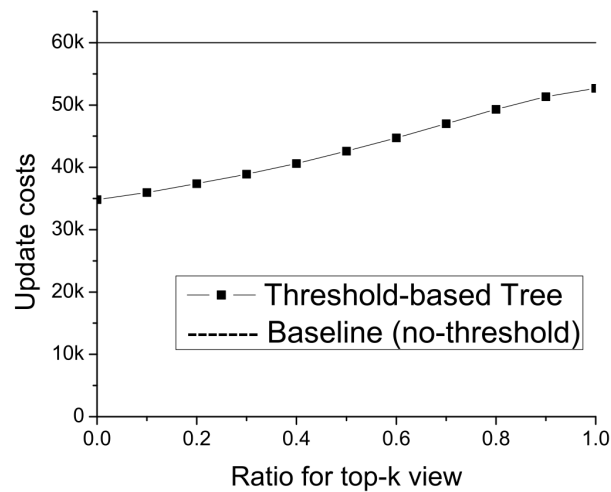


- Synthetic dataset: triplet $\langle A, B, C \rangle$
 - A, is randomly picked from 50 distinct tags.
 - B is numeric, randomly distributed in $[0, 32]$.
 - C is bounded by 50, following uniform distribution and Zipf distribution.
- Two data batches:
 - Loading: populating the data store and initializing thresholds
 - Performance evaluation
- Platform setup:
 - Software: Implemented on top of HBase
 - Hardware: Twenty commodity machines.

Preliminary results: Update costs



**SYRACUSE
UNIVERSITY**
**ENGINEERING
& COMPUTER
SCIENCE**



Summary



- We study the problem of interactive top- k aggregation query over dynamic data.
- We propose KTV-TREE, which combines the threshold based mechanism with materialized views
- KTV tree achieves the fast top- k aggregation processing with reasonably degraded accuracy.
- Future work includes more mature prototyping of KTV-TREE (e.g. on Spark) and experimentation.

Questions?



**SYRACUSE
UNIVERSITY**
**ENGINEERING
& COMPUTER
SCIENCE**

Thank you

Contact:

Yuzhe Tang

Syracuse University

Email: ytang100@syr.edu

Web: ecs.syr.edu/faculty/yuzhe

