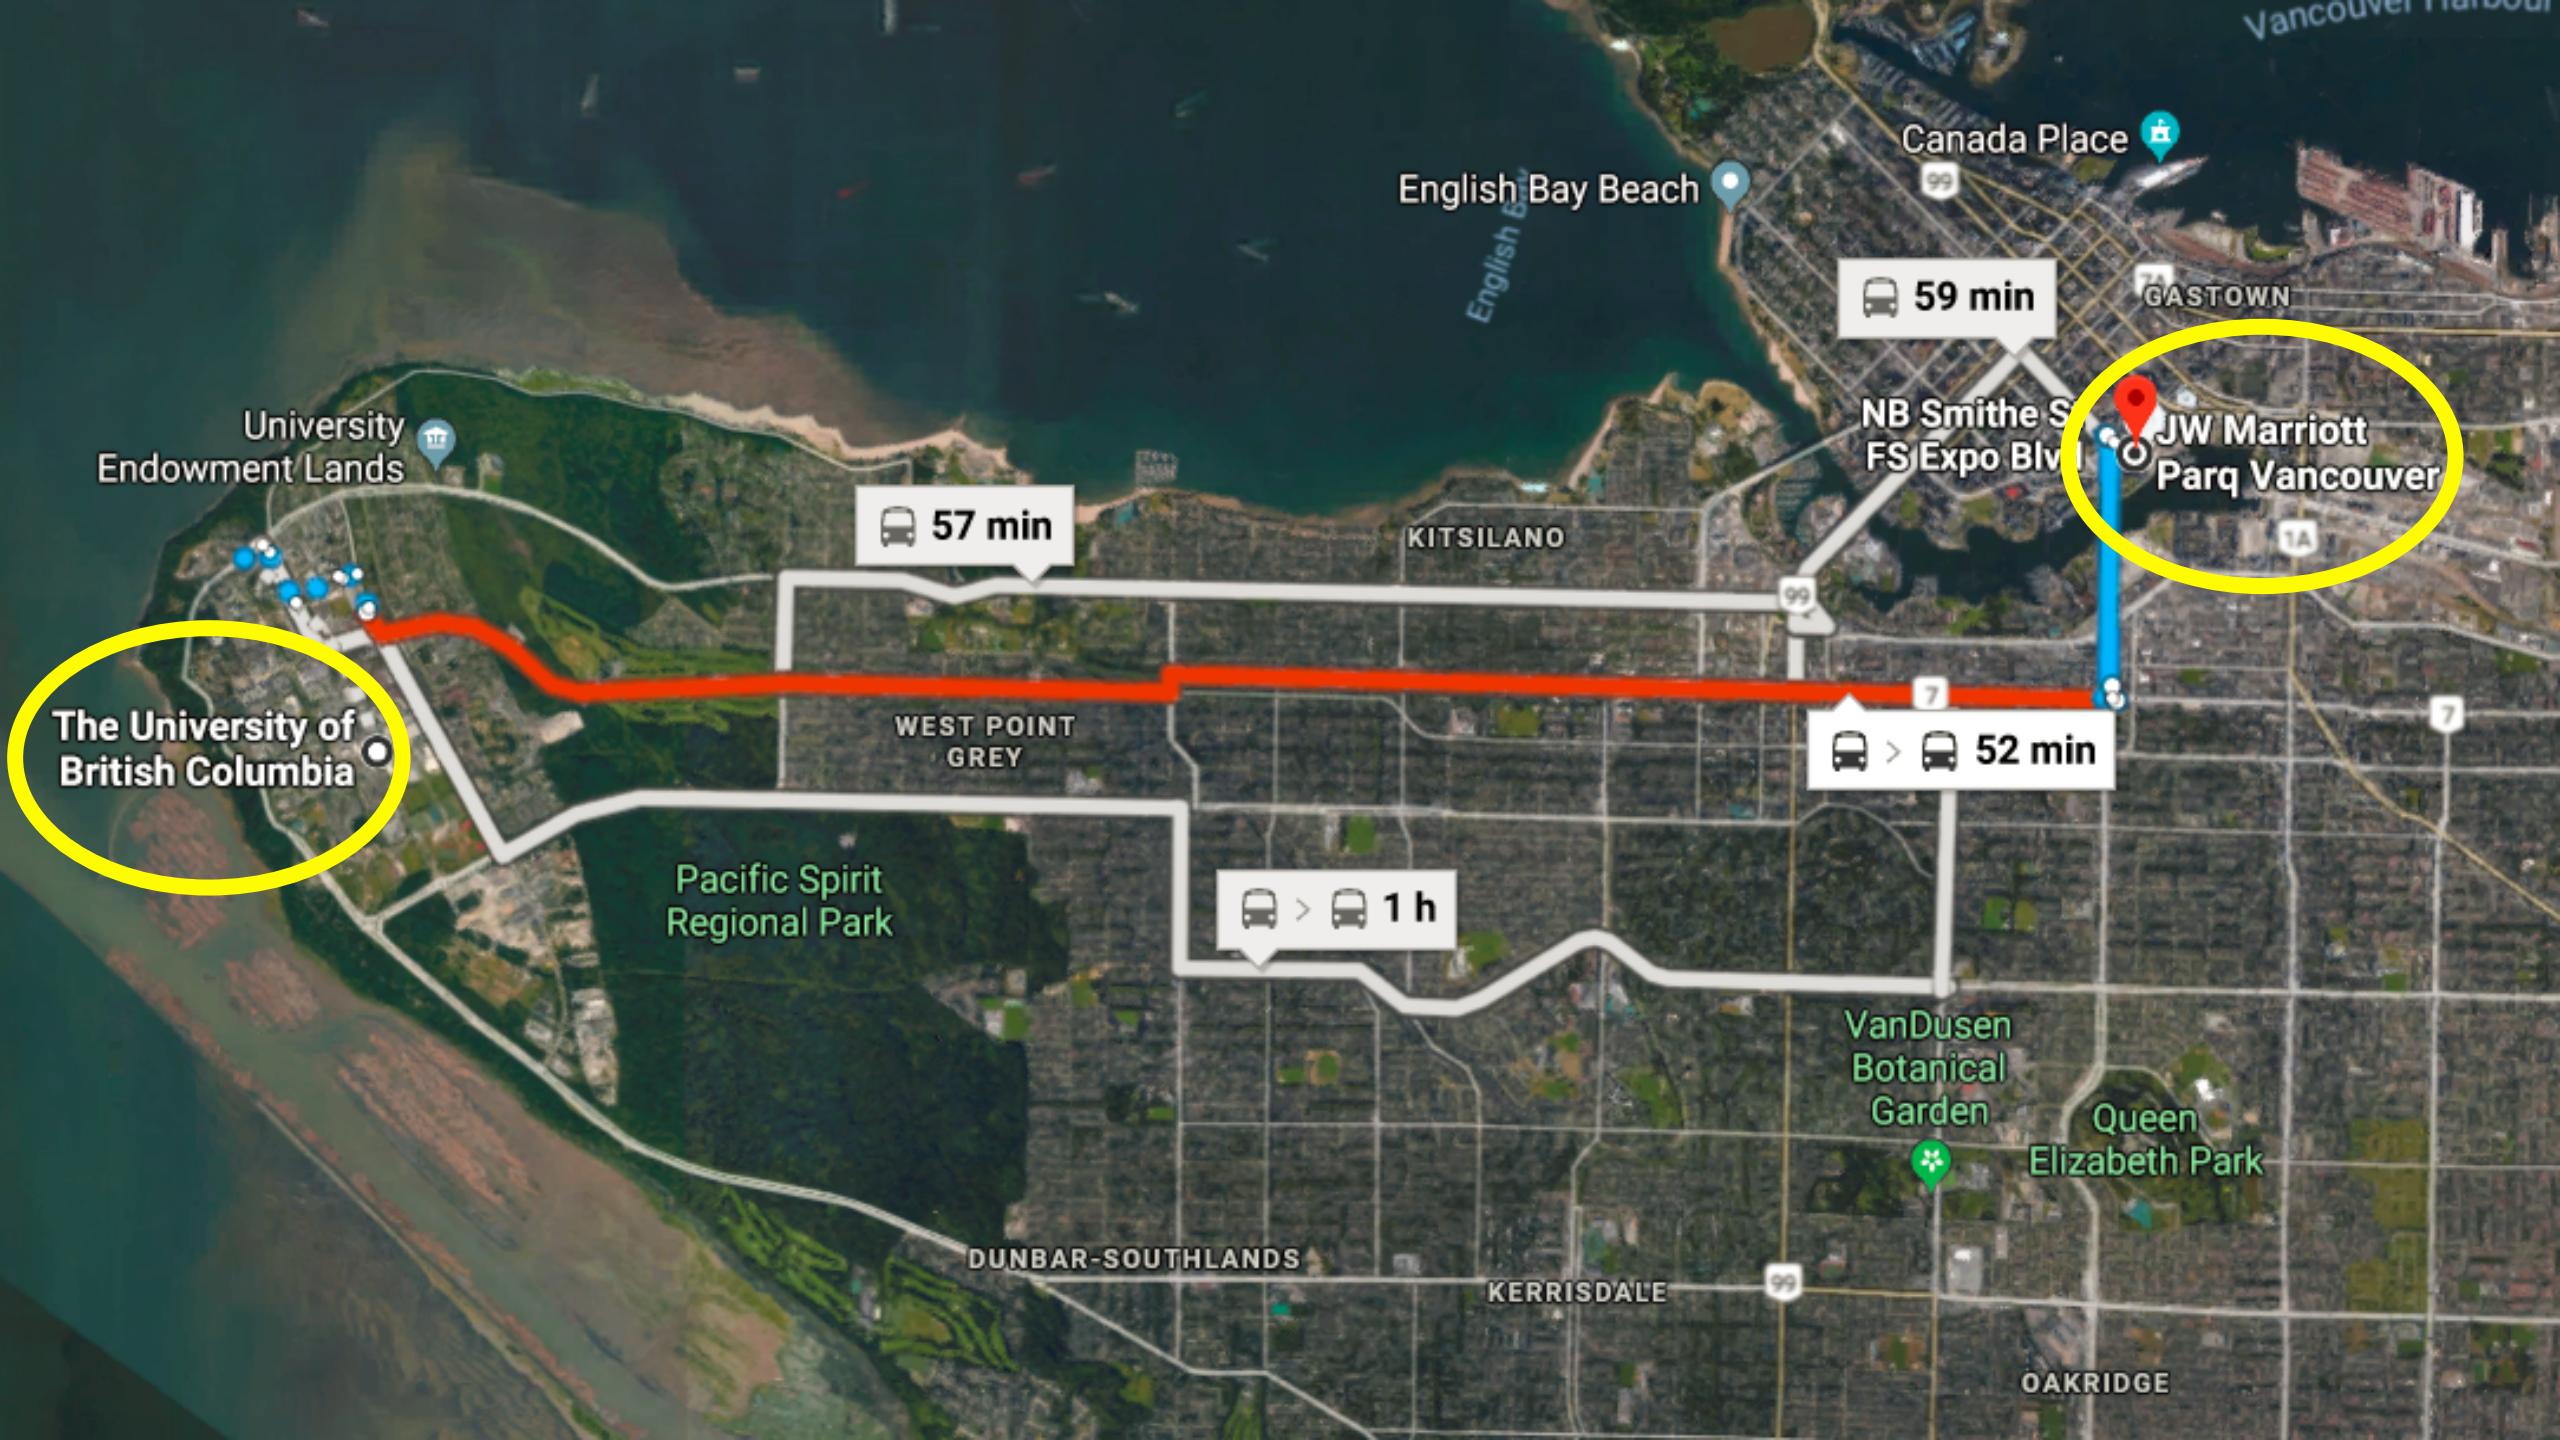
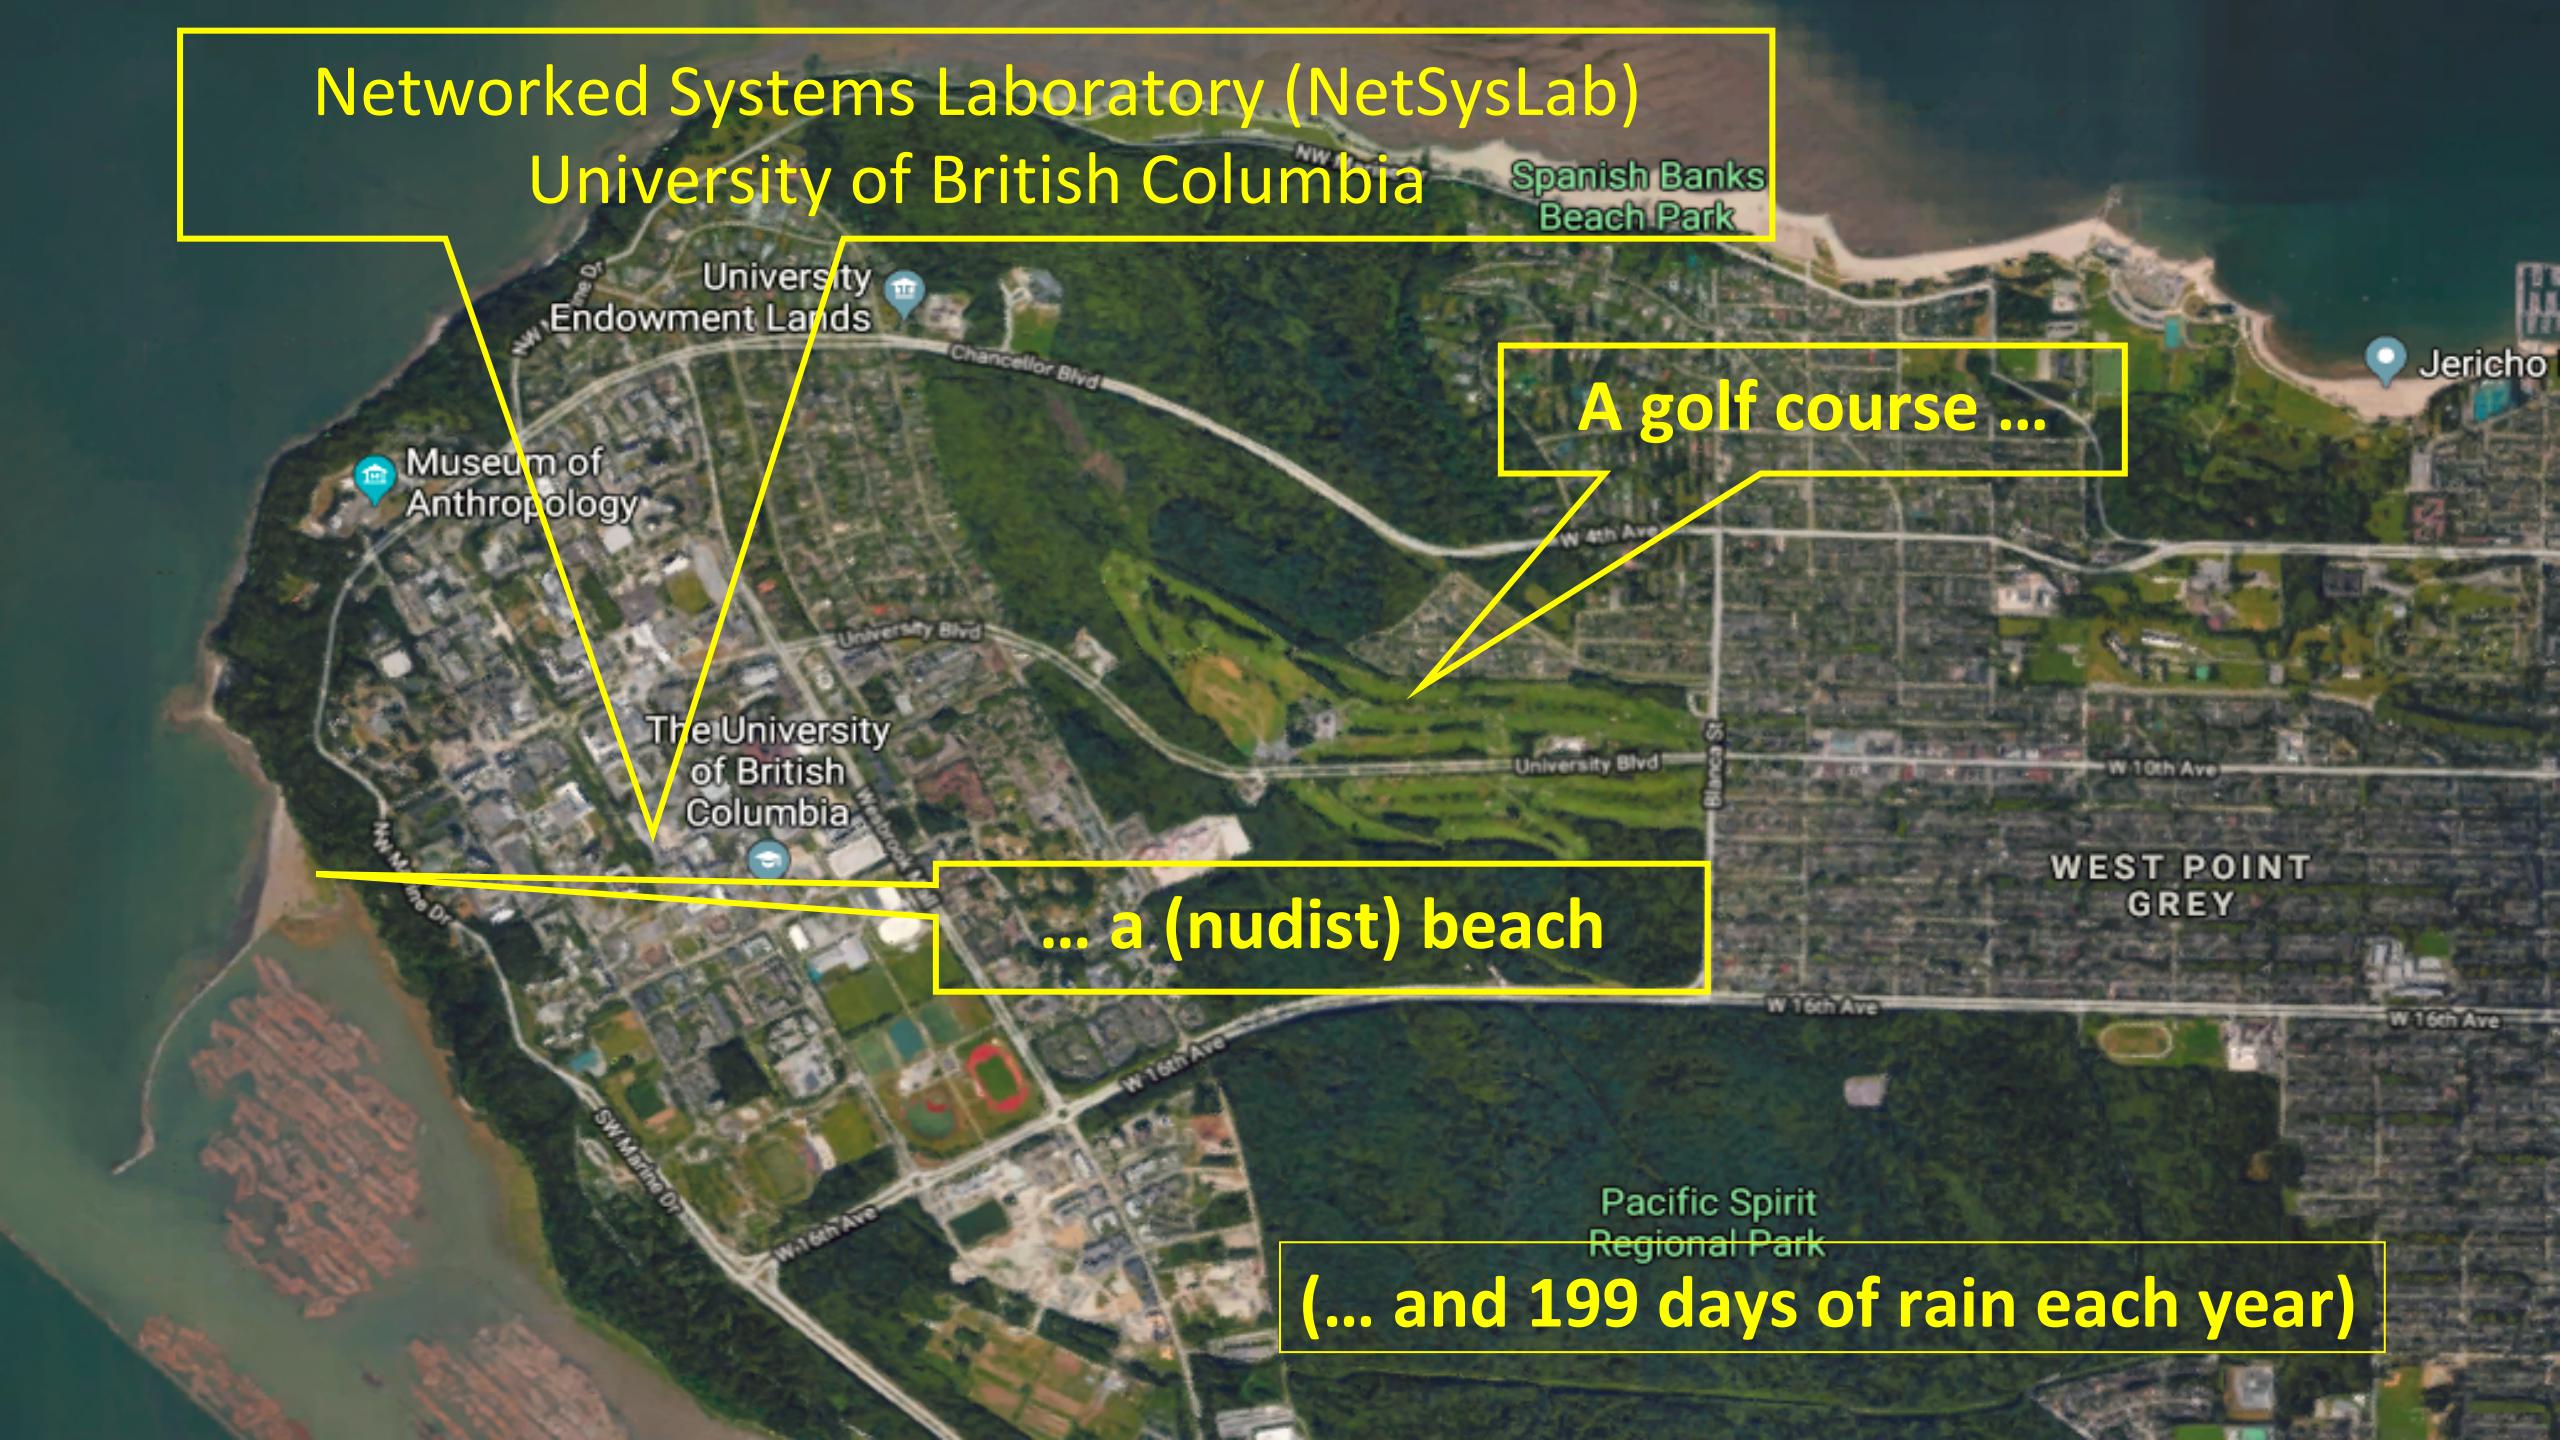


# How well do CPU, GPU and Hybrid Graph Processing Frameworks Perform?

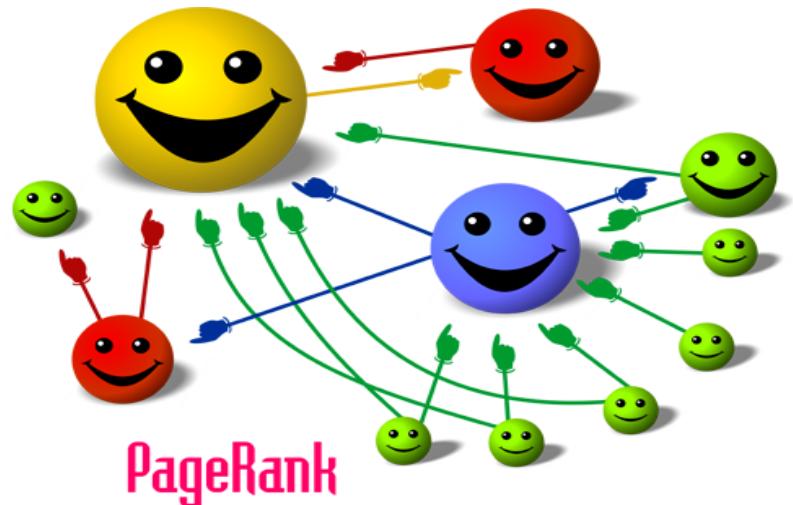
Tanuj Kr Aasawat, Tahsin Reza, Matei Ripeanu  
Networked Systems Laboratory ([NetSysLab](#))  
University of British Columbia



# Networked Systems Laboratory (NetSysLab) University of British Columbia



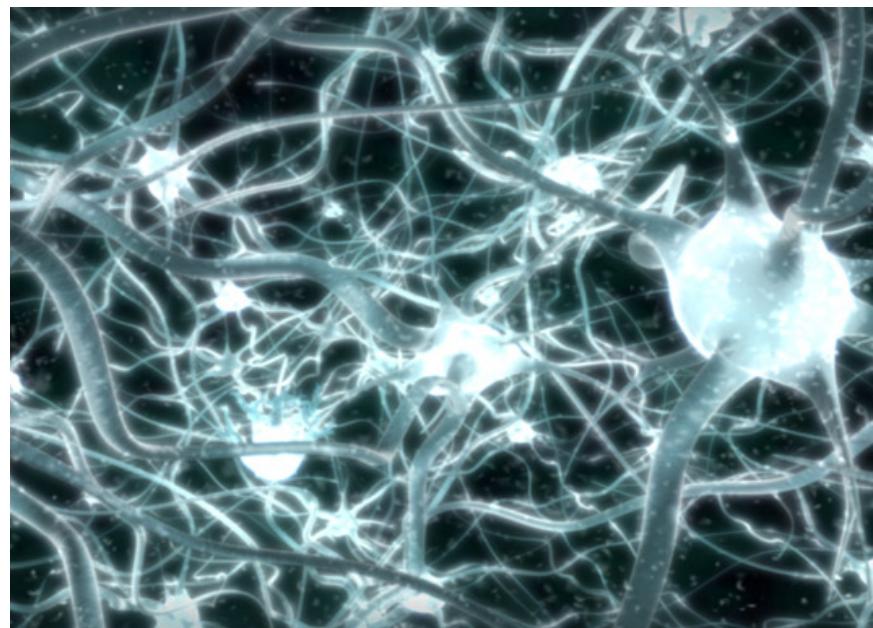
# Graphs are Everywhere



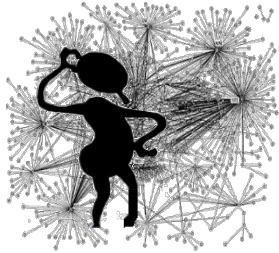
1B users  
150B friendships



100B neurons  
700T connections



# Challenges in Graph Processing



*Poor locality*

*Data-dependent memory  
access patterns*

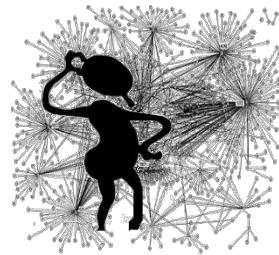
*Low compute-to-  
memory access ratio*

**Graph500 “mini” graph  
requires 128 GB.**

*Large memory footprint*

*Varying degrees of parallelism  
(both intra- and inter- stage)*

# Processing Elements Characteristics



Graph500 “mini” graph  
requires 128 GB.

*Poor locality*  
*Data-dependent memory access patterns*

*Low compute-to-memory access ratio*

*Large memory footprint*

*Varying degrees of parallelism  
(both intra- and inter- stage)*

## CPUs



Large Caches

## GPUs



Caches

Massive hardware multithreading

>1TB

~16GB

*Assemble a hybrid platform?*

# Graph Processing Frameworks

Programming Model  
(Vertex Programming/Linear Algebra)



**High Performance**

Architecture  
(Single-node or ~~Distributed~~)

**CPU/GPU/Hybrid**

# Motivation

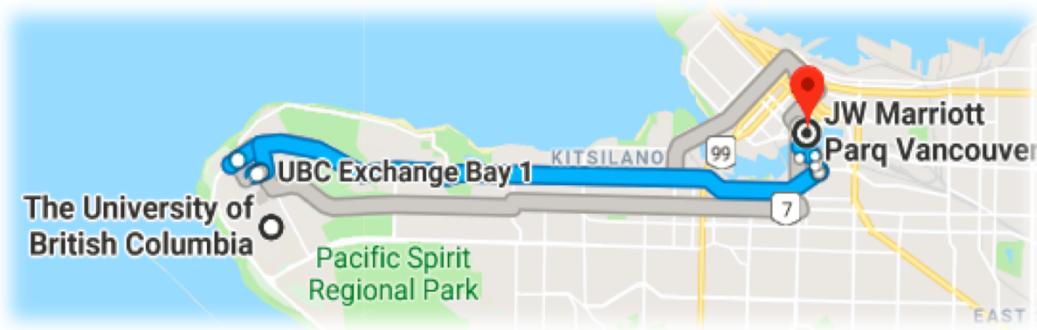
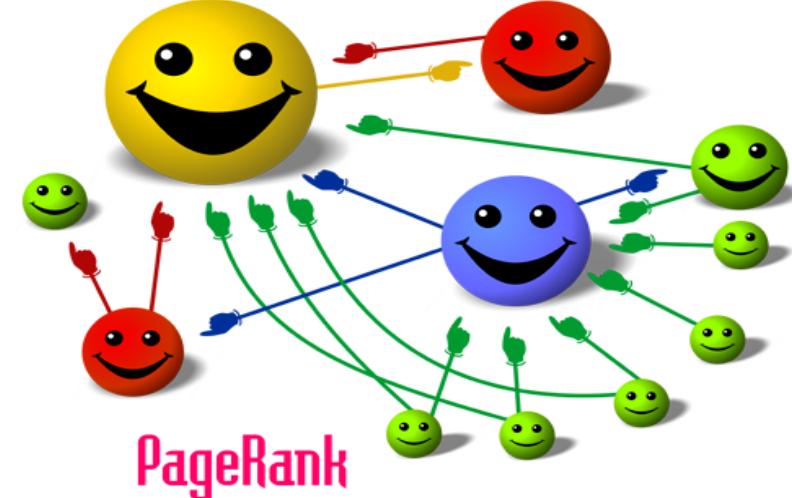
How architecture and programming model combination improves **performance** and **efficiency** of the system as a whole?

# Graph Processing Frameworks

Framework	Architecture	Programming Model
Galois <a href="#">UTexas, Austin</a>	CPU	Vertex Programming
GraphMat <a href="#">Intel</a>	CPU + Distributed	Linear Algebra
Gunrock <a href="#">UC, Davis</a>	Multi - GPU	Vertex Programming
Nvgraph <a href="#">Nvidia</a>	GPU	Linear Algebra
Totem <a href="#">UBC</a>	CPU + multi-GPU	Vertex Programming

# Benchmark Algorithms

- PageRank
  - Ranking web pages
  - Compute intensive
- Single Source Shortest Paths (SSSP)
  - IP routing, Transportation networks
- Breadth-First Search (BFS)
  - Finding connected component, subroutine
  - Memory intensive



# Evaluation Metrics

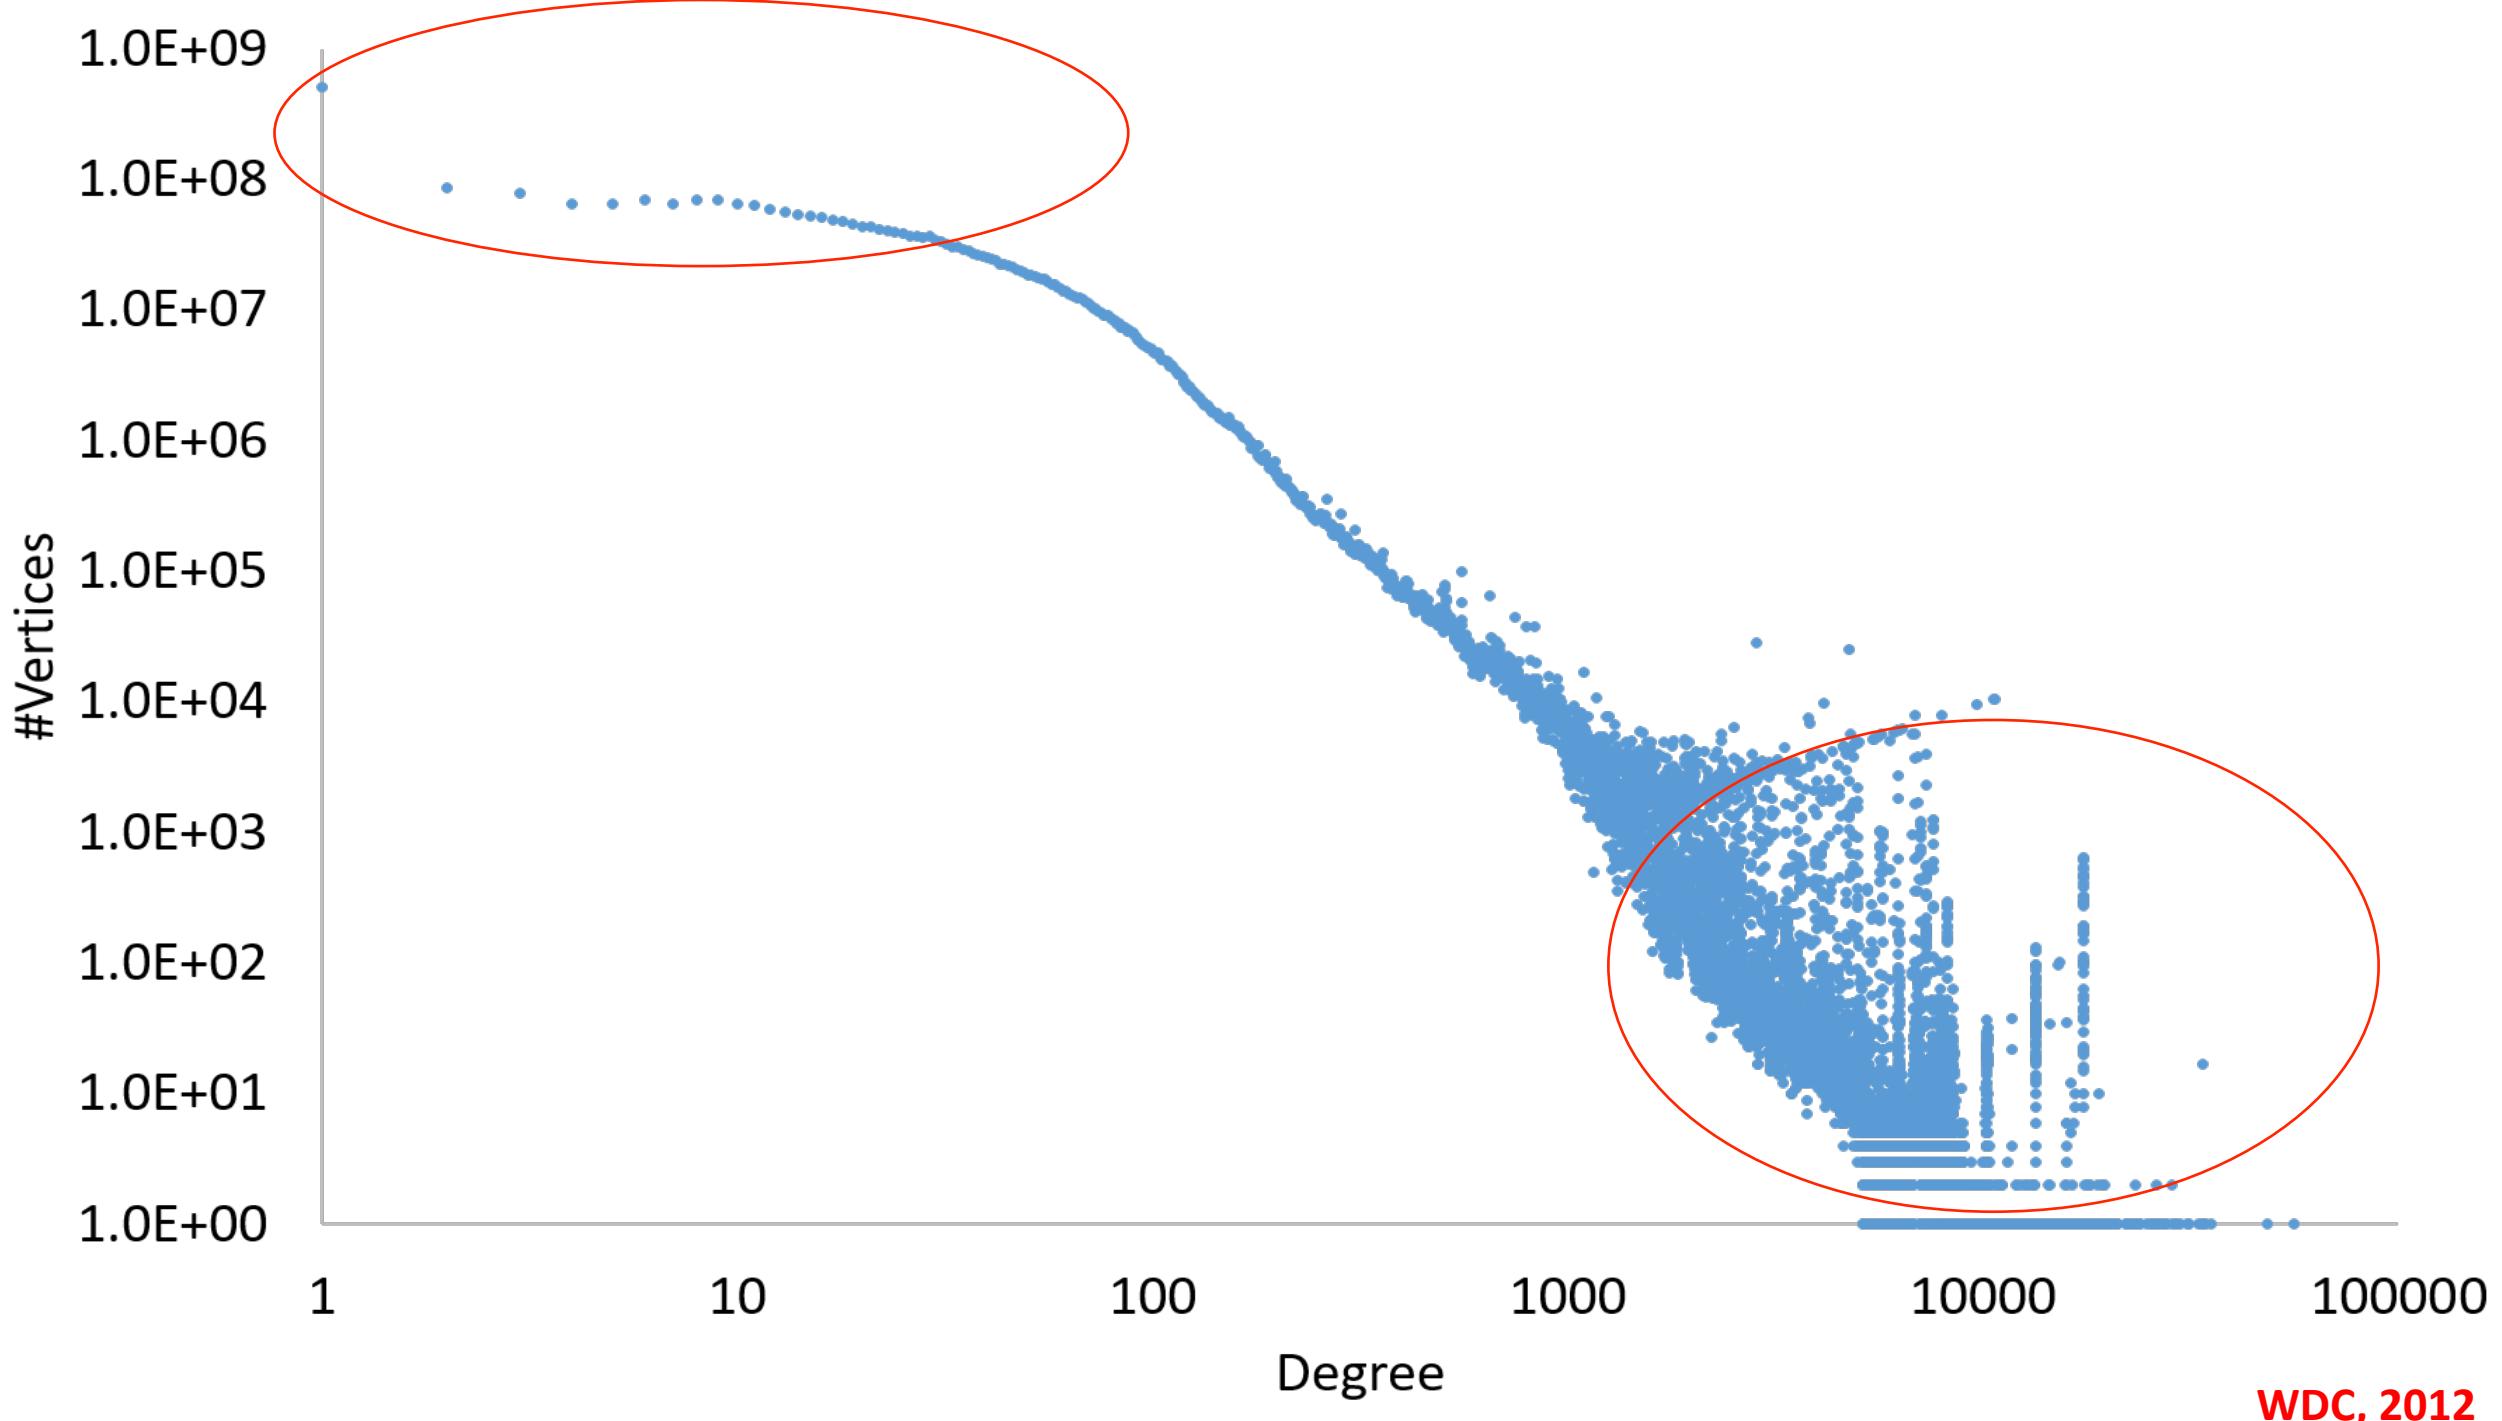
- Raw Performance
  - Traversed Edges Per Second (TEPS): *Traversed Edges / Execution Time*
- Energy Consumption
  - *Average Power consumed \* Execution Time*
- Scalability
  - Strong scaling w.r.t processing units

# Testbed Characteristics

	<b>System 1</b>
CPU	2x Intel Xeon E5-2695 v3 (Haswell)
#CPU Cores	28
Host Memory	<b>512 GB DDR4</b>
L3 Cache	70 MB
PCIe	3.0 – x16
GPU	2x Nvidia Tesla K40c
GPU Thread Count	2880
GPU Memory	<b>12 GB</b>

# Datasets

	Graph	#Vertices	#Edges	Max Degree	Avg. Degree
Real World	Com-Orkut	3 M	234 M	33,313	78
	liveJournal	4.8 M	68 M	20,292	14
	Road-USA	28.8 M	47.9 M	9	1.6
	Twitter	52 M	3.9 B	3,691,240	75
Synthetic	RMAT22	4 M	128 M	168,729	32
	RMAT23	8 M	256 M	272,808	32
	RMAT24	16 M	512 M	439,994	32
	RMAT27	128 M	4 B	3,910,241	32



WDC, 2012

# Memory Consumption

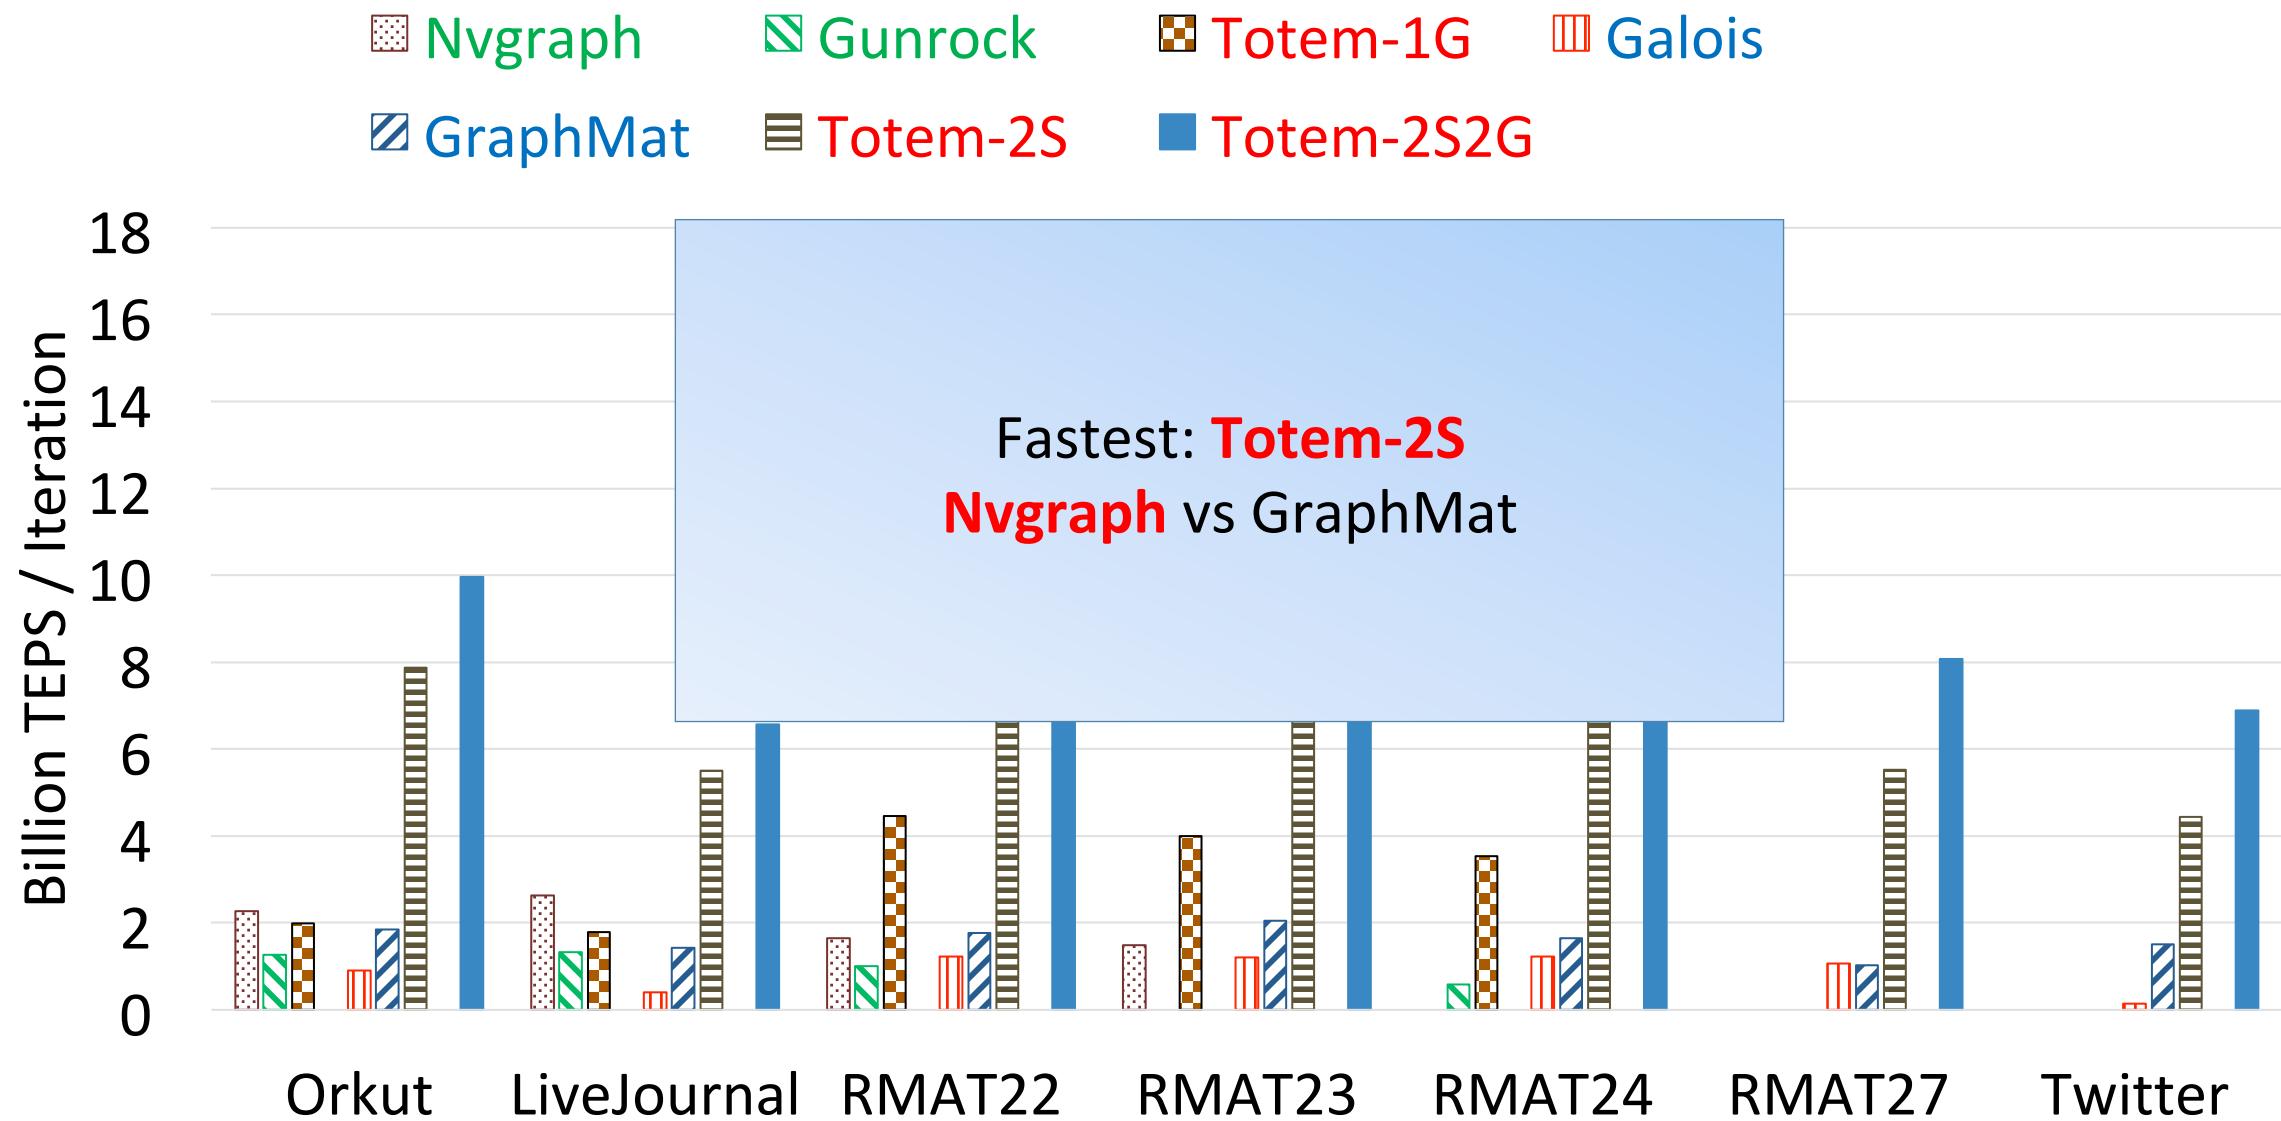
Framework	Memory layout	PageRank	SSSP	BFS
Nvgraph	CSC (PageRank, SSSP) and CSR (BFS)	1,159 (1.8x)	<b>1,111 (1.0x)</b>	<b>683 (1.0x)</b>
Gunrock	CSR and COO	<b>641 (1.0x)</b>	1,582 (1.4x)	1,443 (2.1x)
Galois	CSR	1,599 (2.5x)	2,074 (1.9x)	1,432 (2.1x)
GraphMat*	DCSC	<b>2,818 (4.4x)</b>	<b>2,786 (2.5x)</b>	<b>2,980 (4.4x)</b>
Totem-2S	CSR	1,275 (2.0x)	2,198 (2.0x)	1,282 (1.9x)
Totem-2S2G	CSR	1,628 (2.5x)	2,587 (2.3x)	1,658 (2.4x)

9,354 MB  
during pre-  
processing  
step

Memory Consumption (in MB) for RMAT22 graph (edge list size: 512 MB)

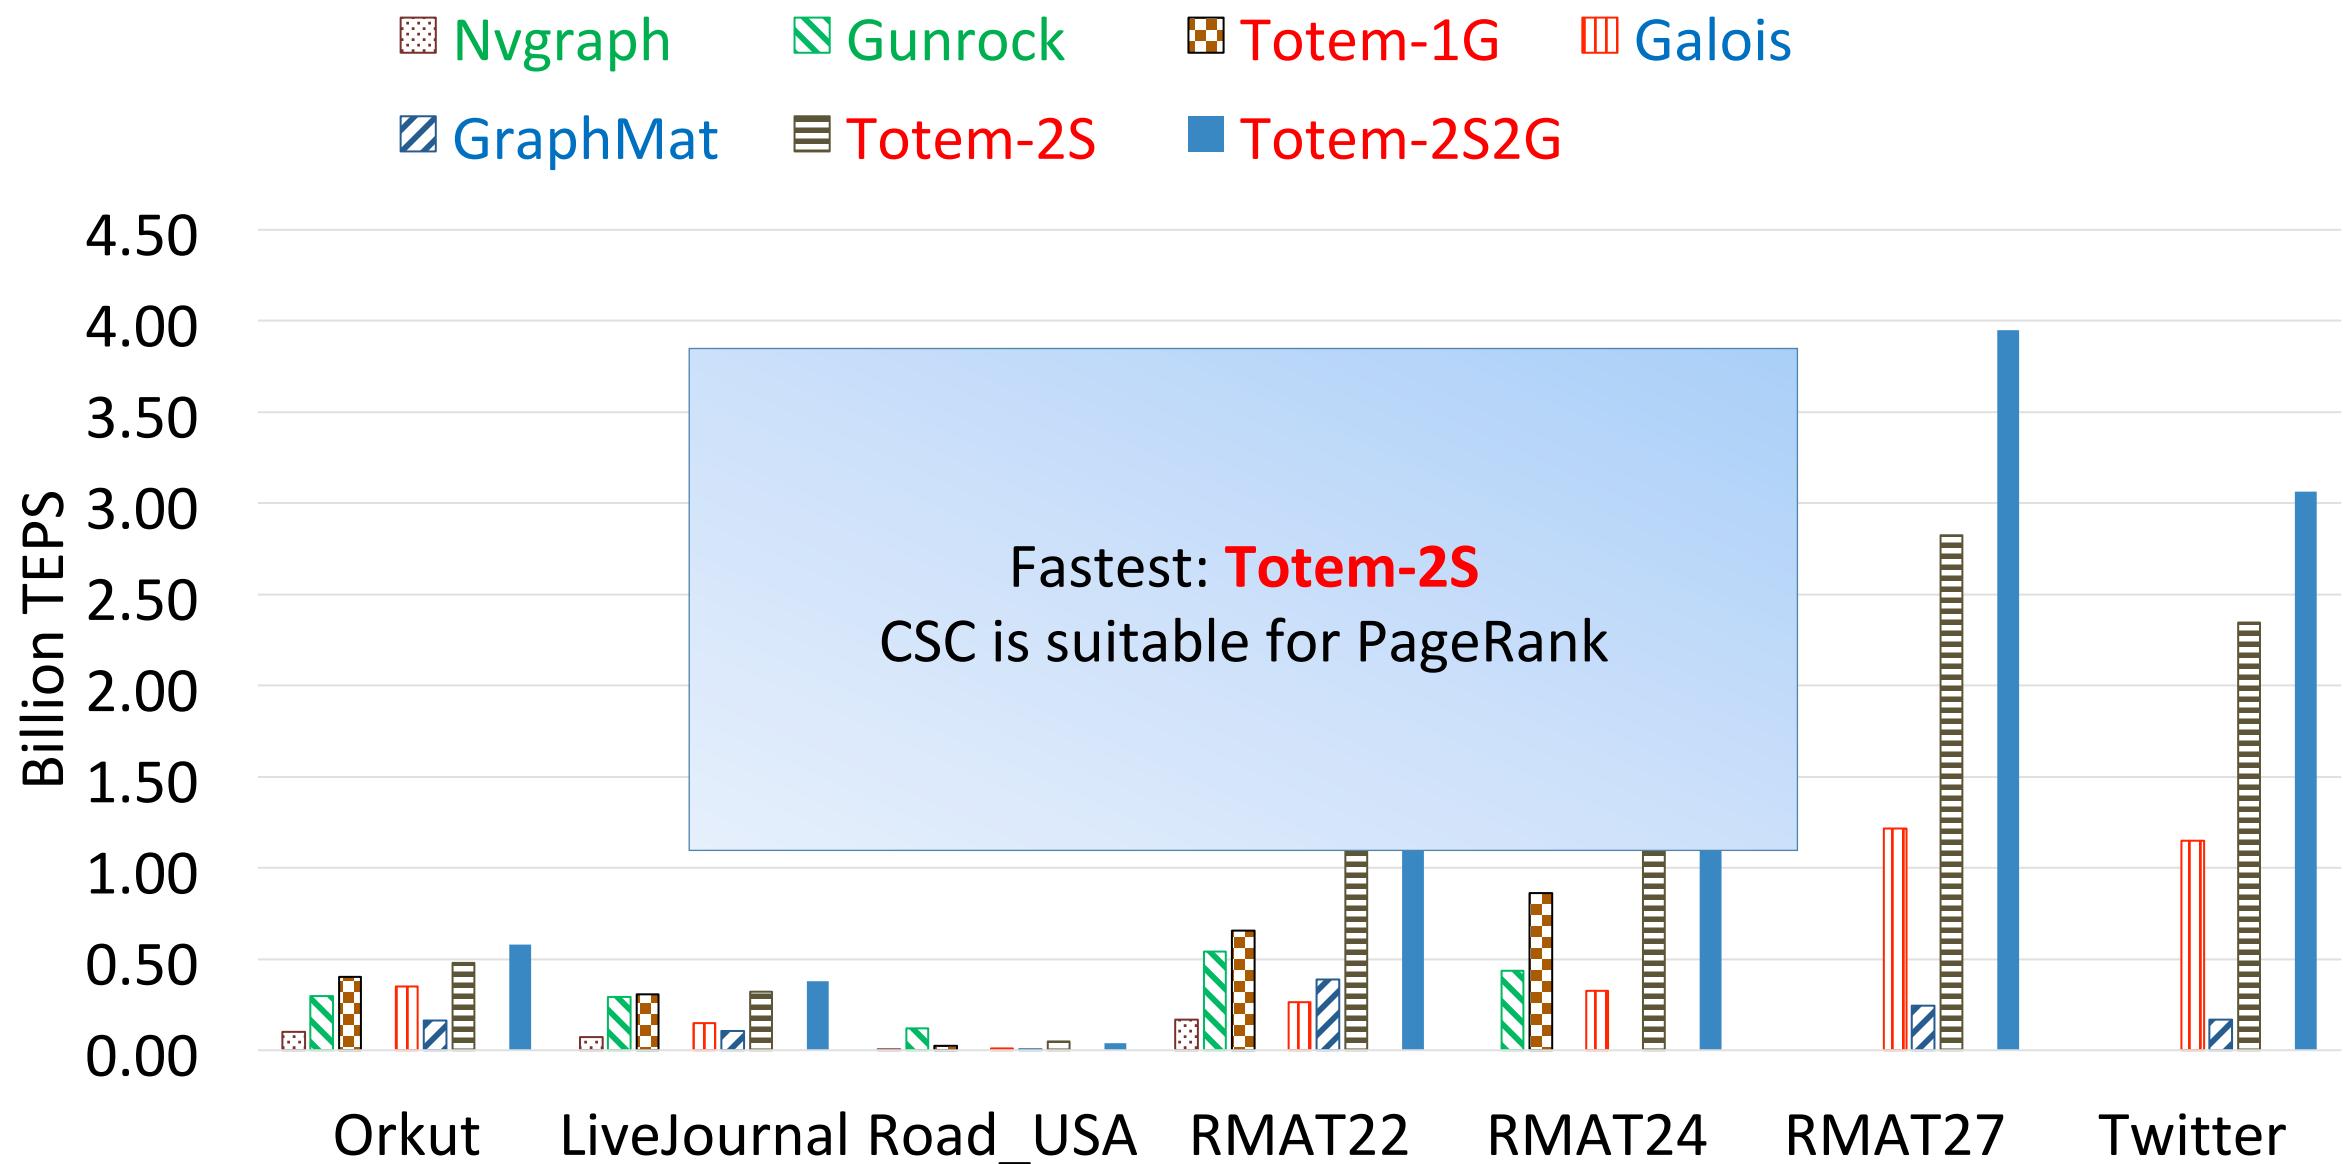
# Experimental Results

## 1. Raw Performance - PageRank

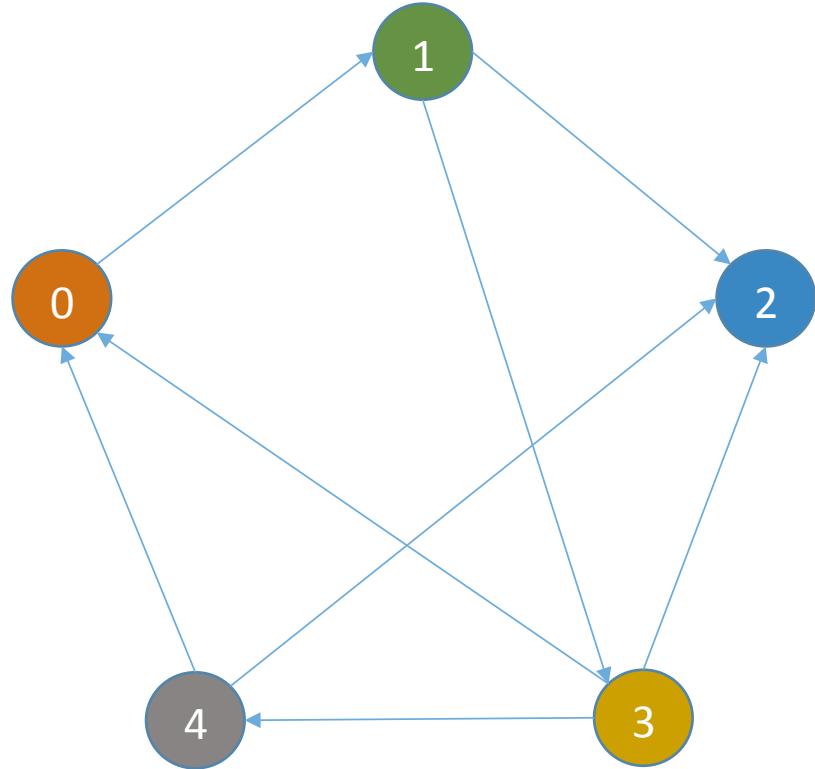


# Experimental Results

## 1. Raw Performance - SSSP



# Graph Layout in Memory



CSR Representation

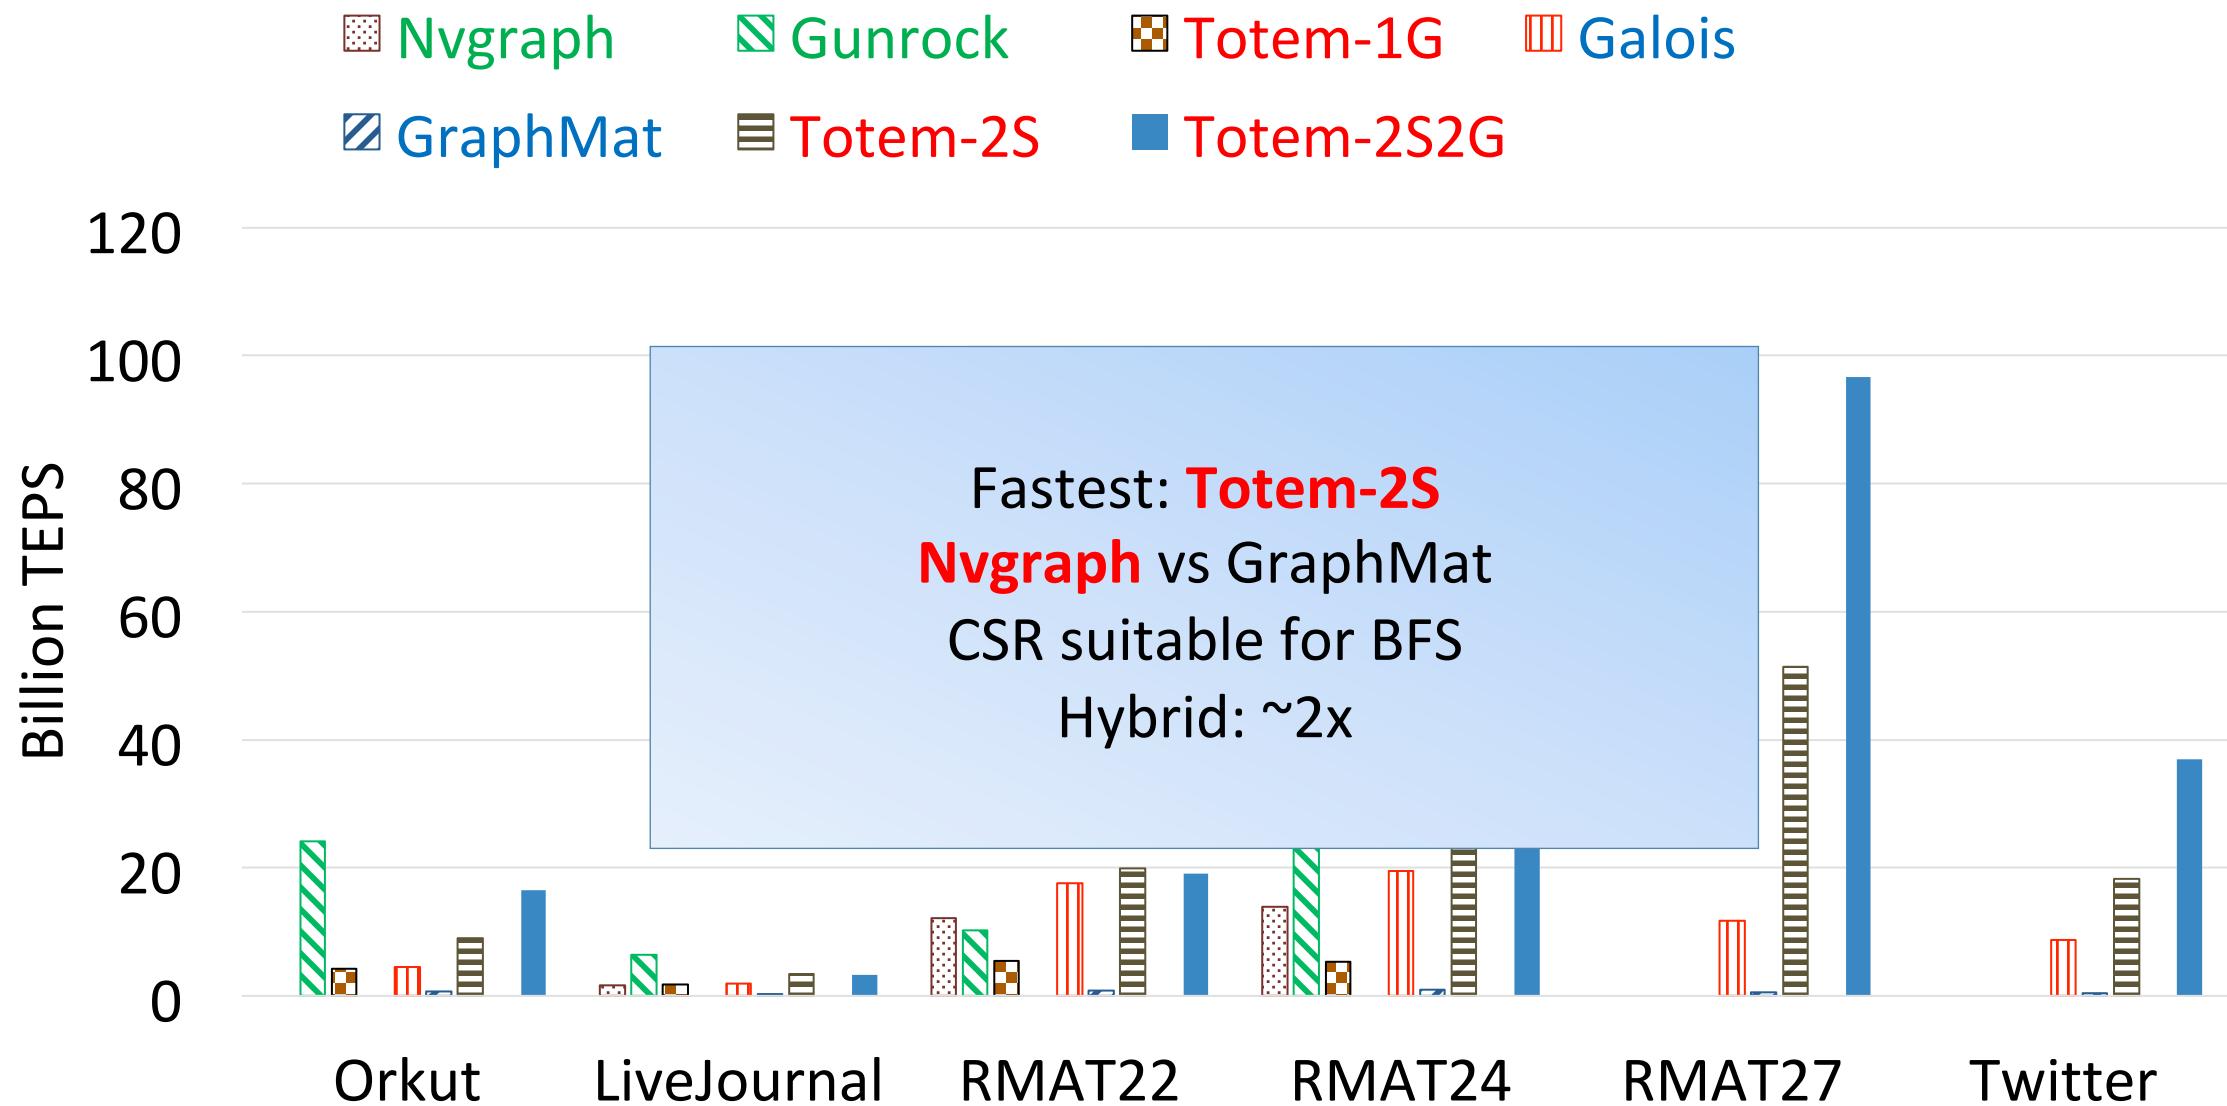
rowPtr	0	1	3	3	6	8
VertexId	0	1	2	3	4	5*
edgeList	1	2	3	0	2	4
	0	1	2	3	4	5

CSC Representation

colPtr	0	2	3	6	7	8
VertexId	0	1	2	3	4	5*
edgeList	3	4	0	1	3	4
	0	1	2	3	4	5

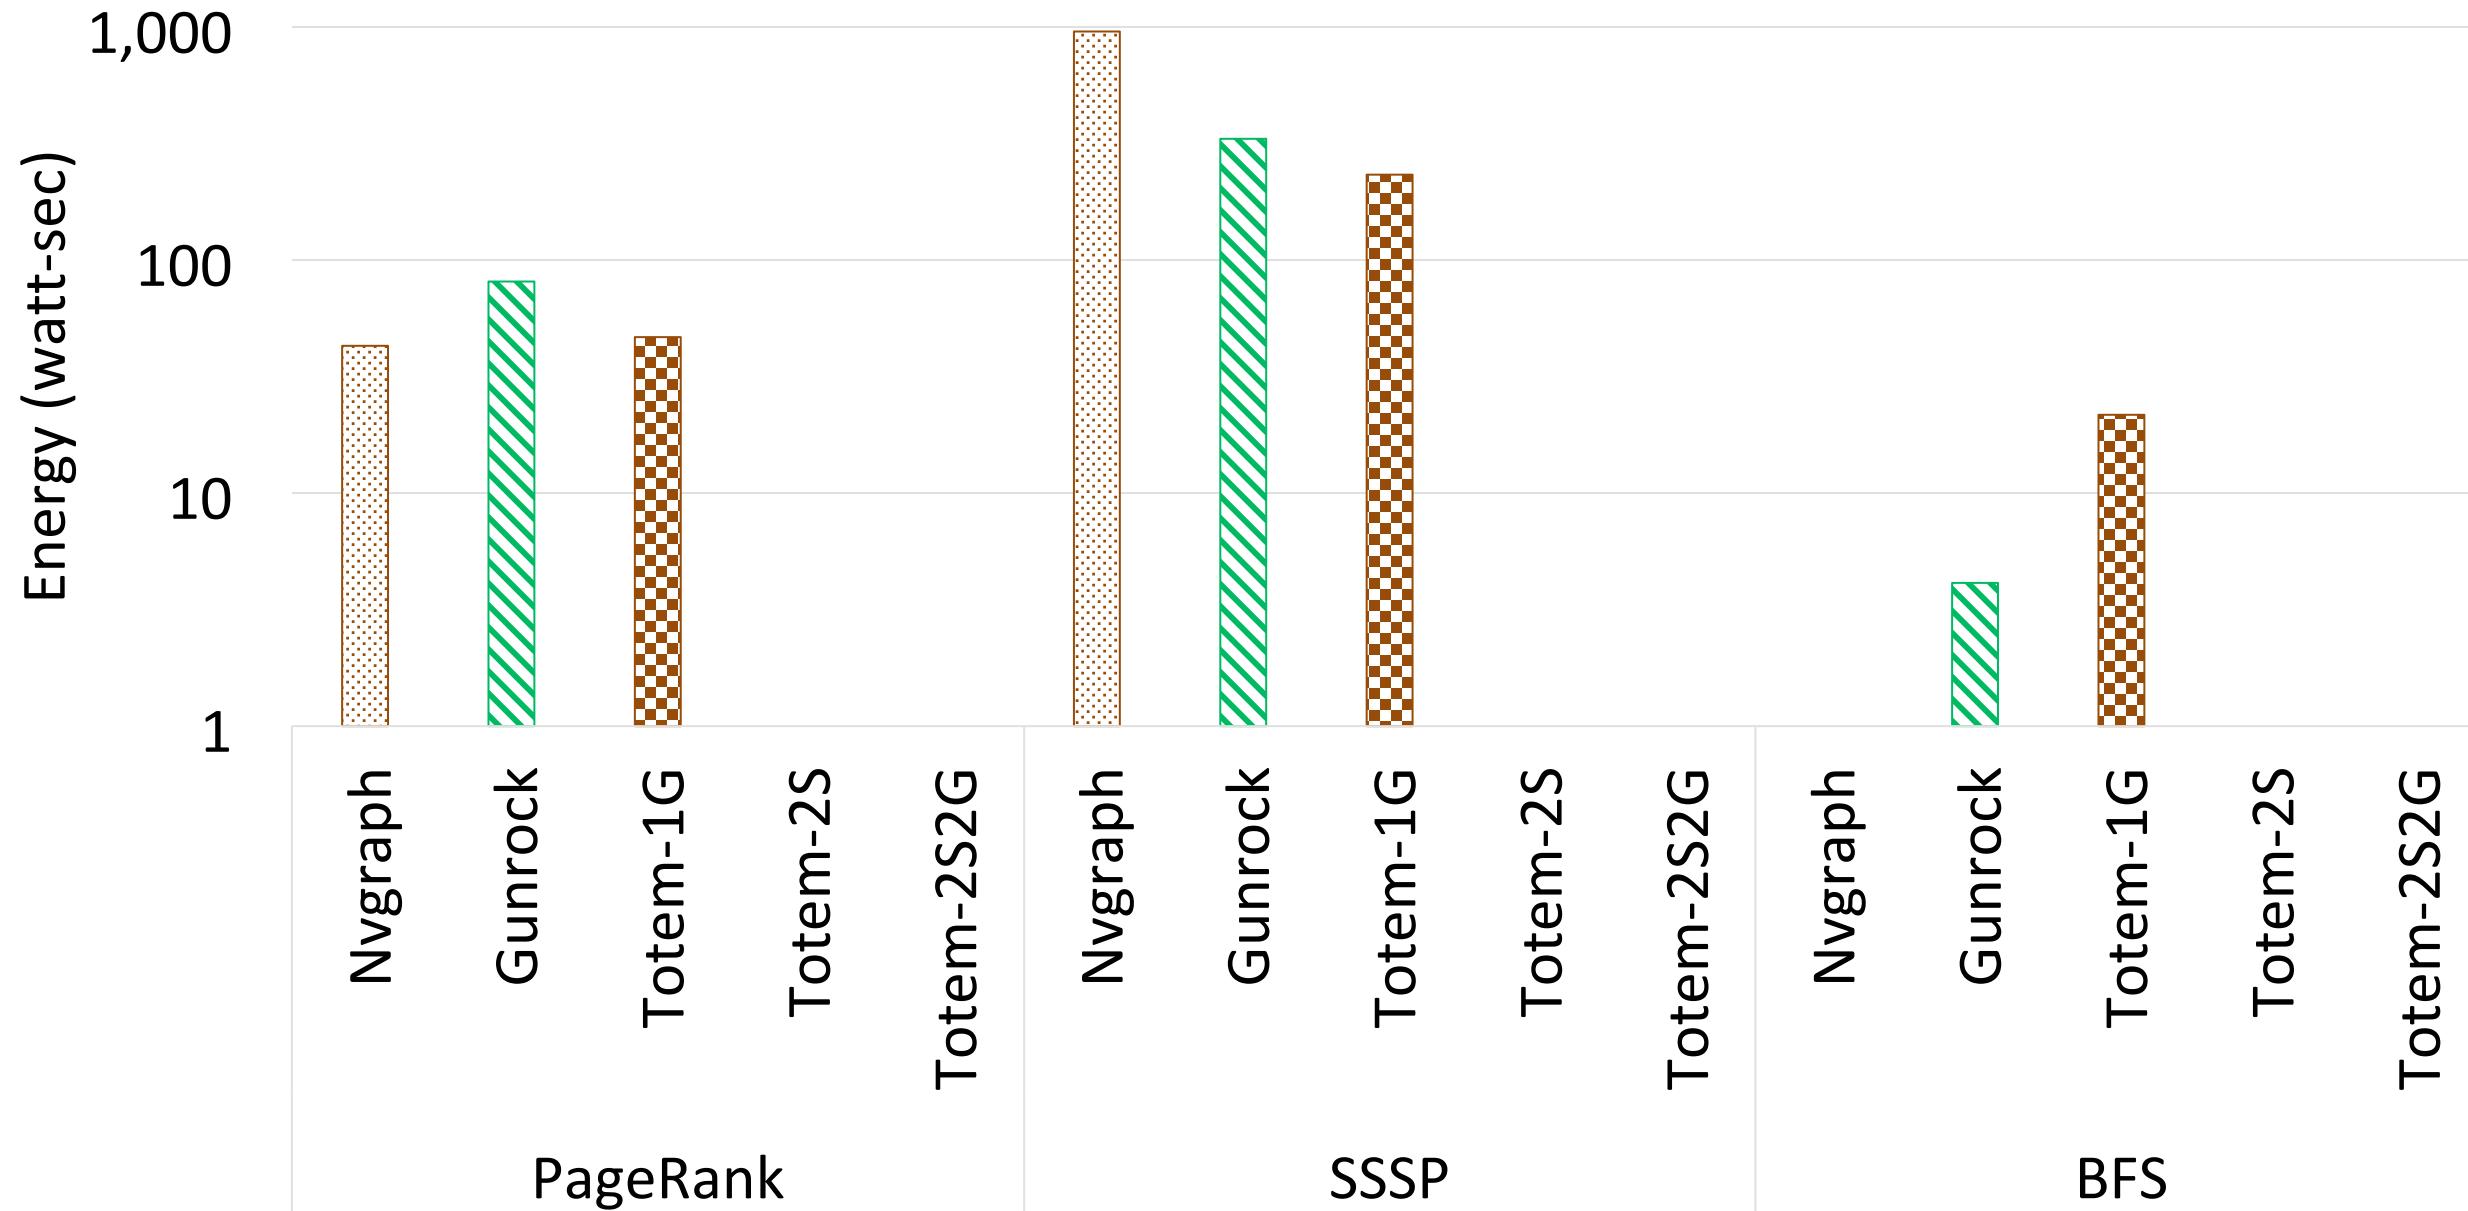
# Experimental Results

## 1. Raw Performance - BFS



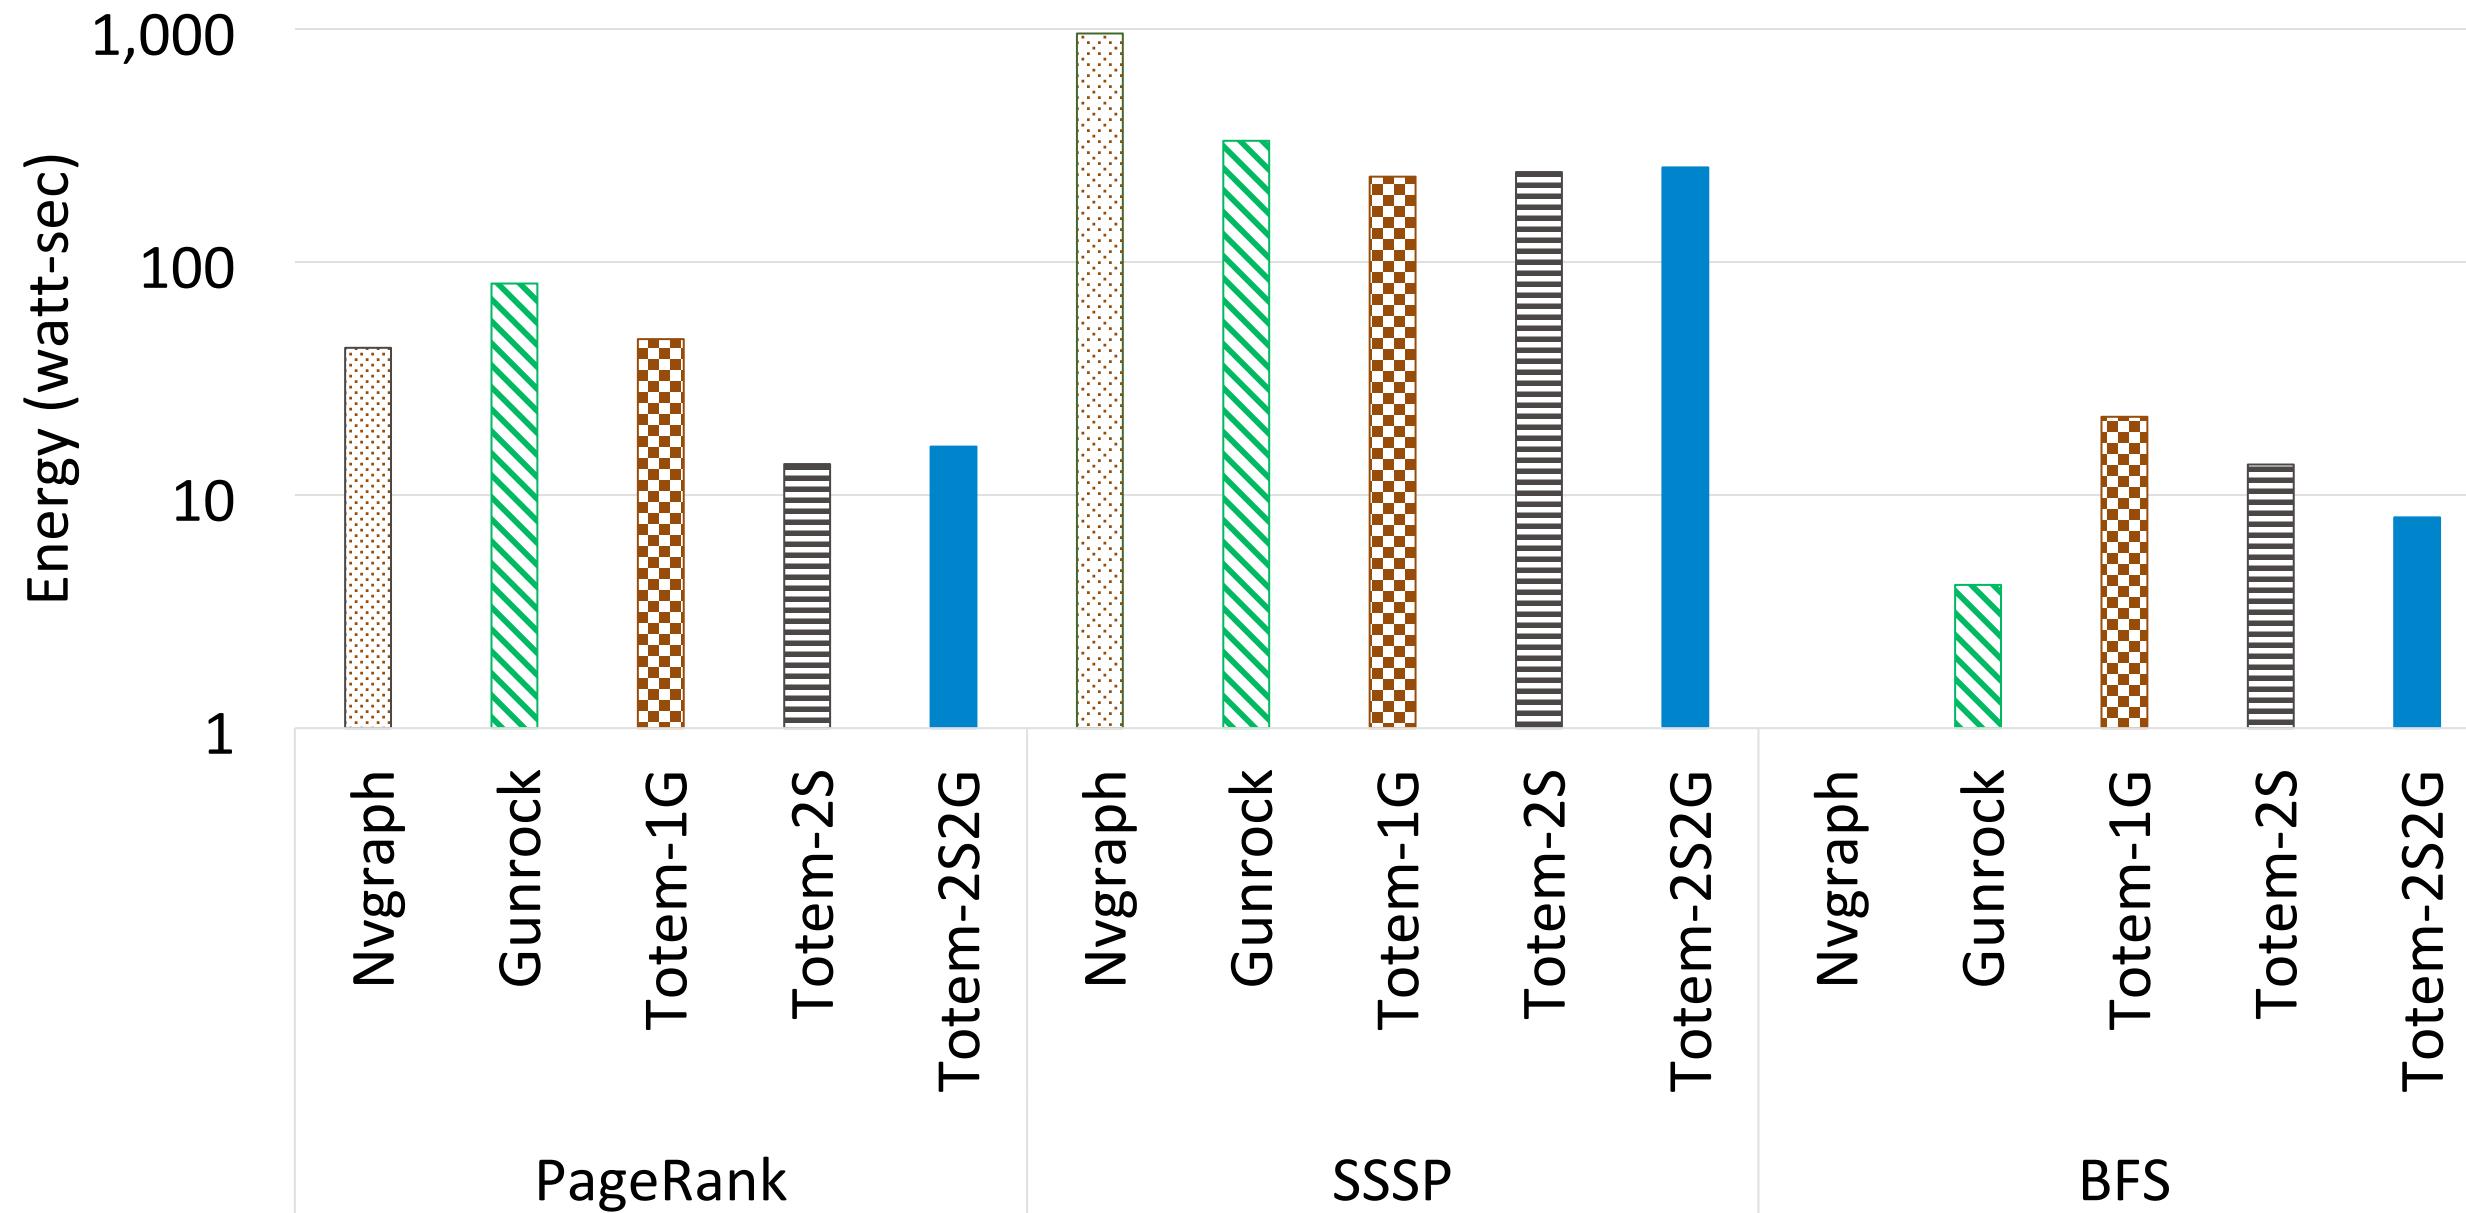
# Experimental Results

## 2. Energy Consumption – GPU Frameworks – Orkut Workload



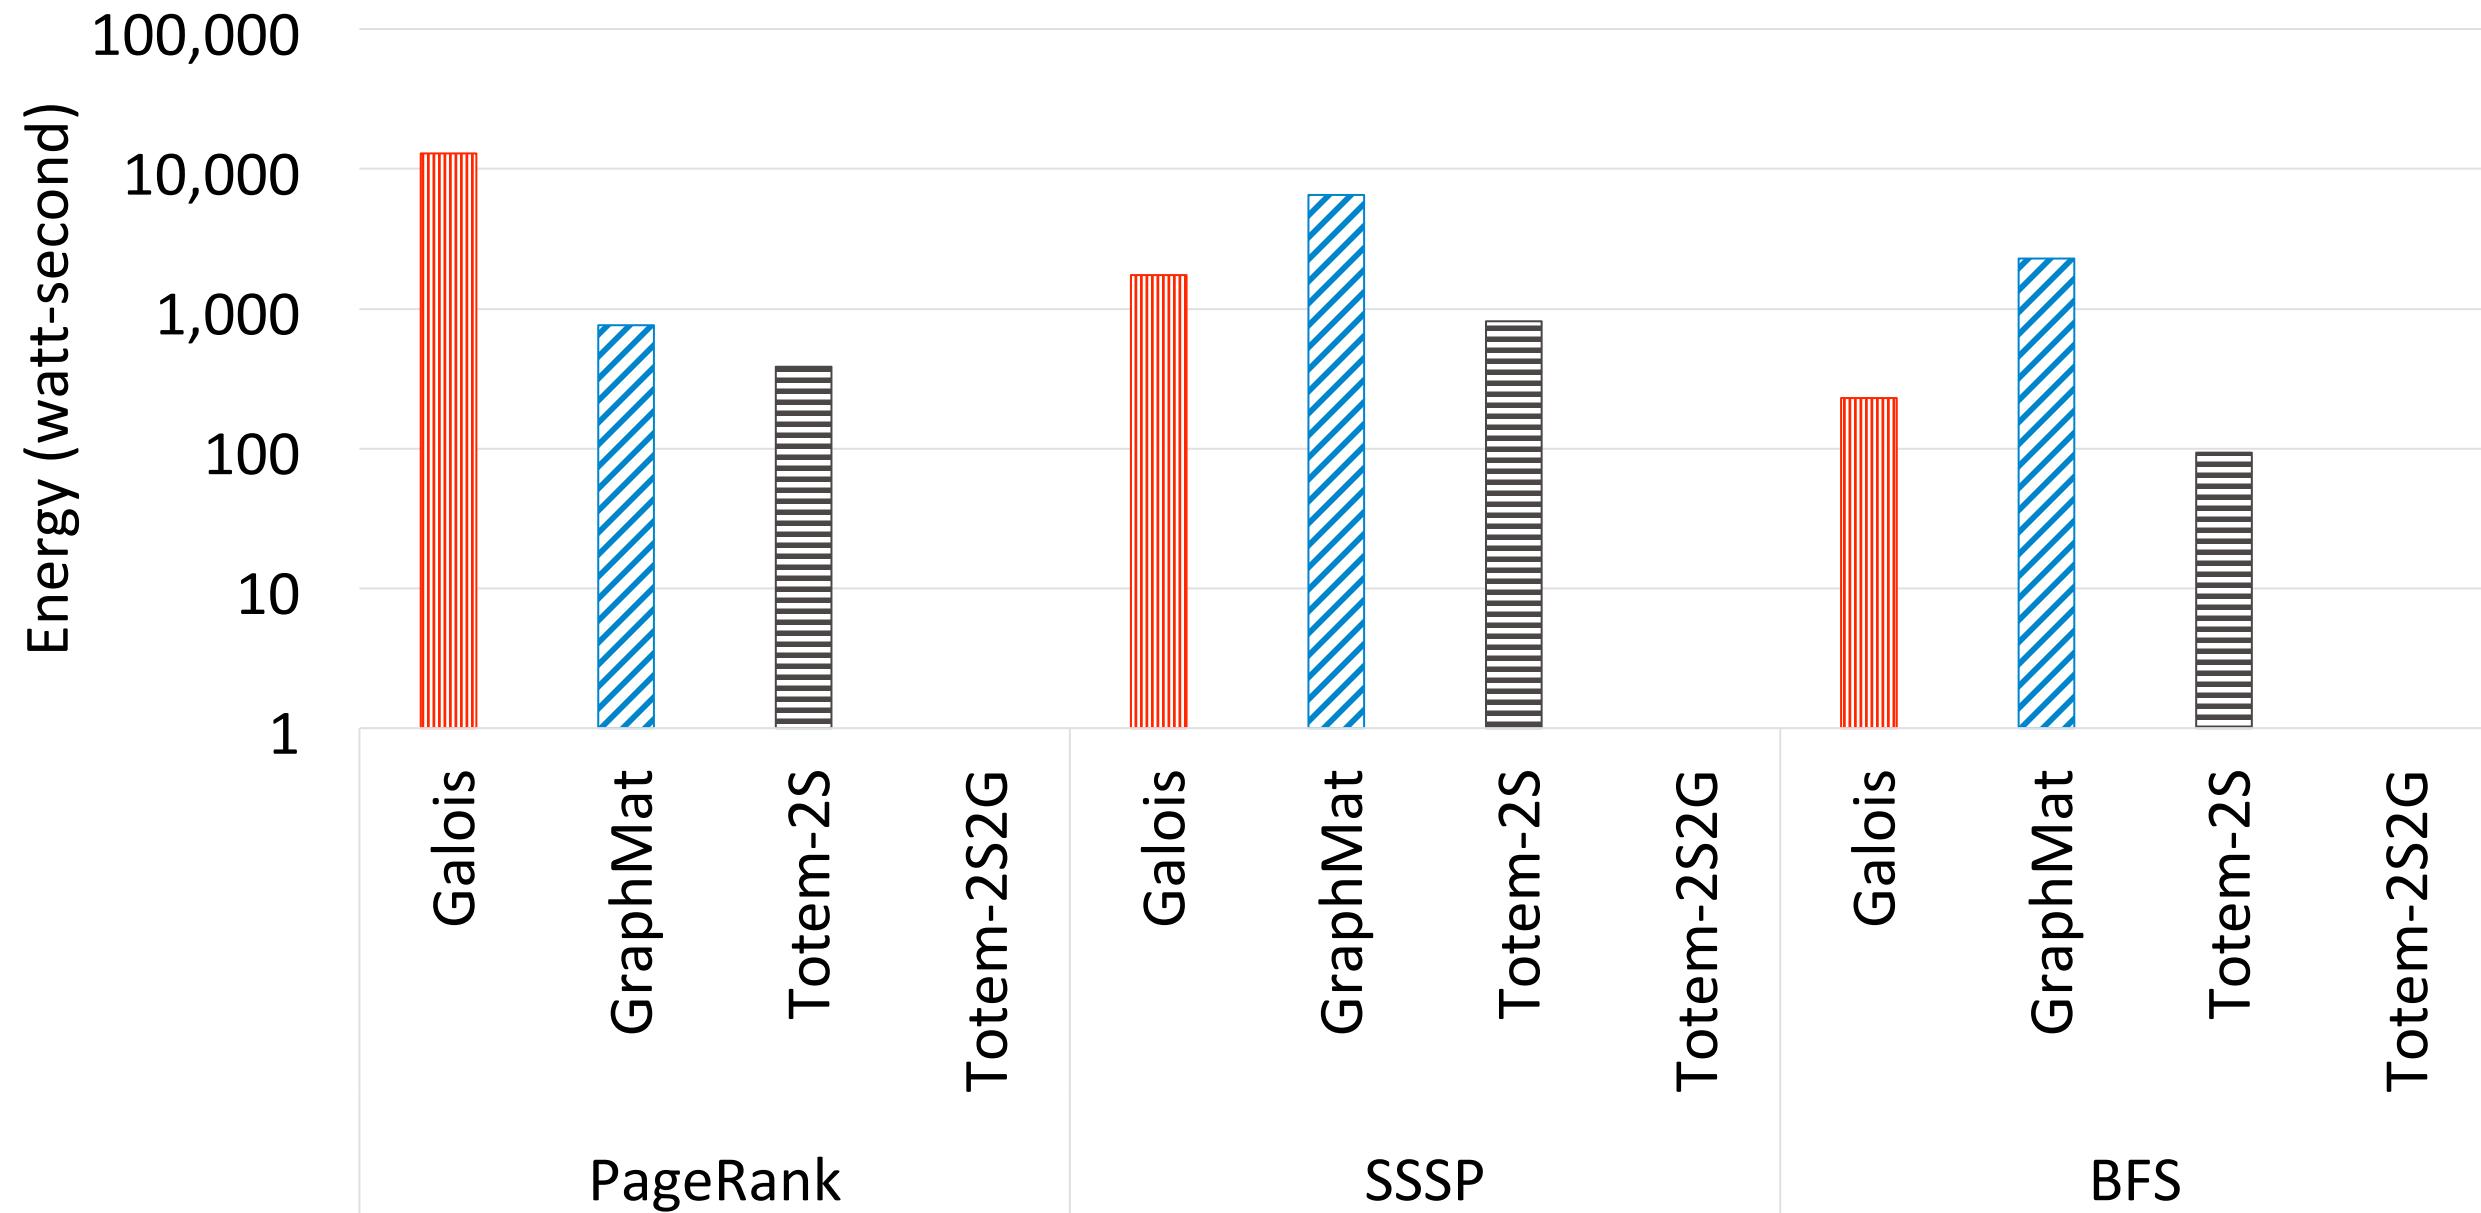
# Experimental Results

## 2. Energy Consumption – GPU Frameworks – Orkut Workload



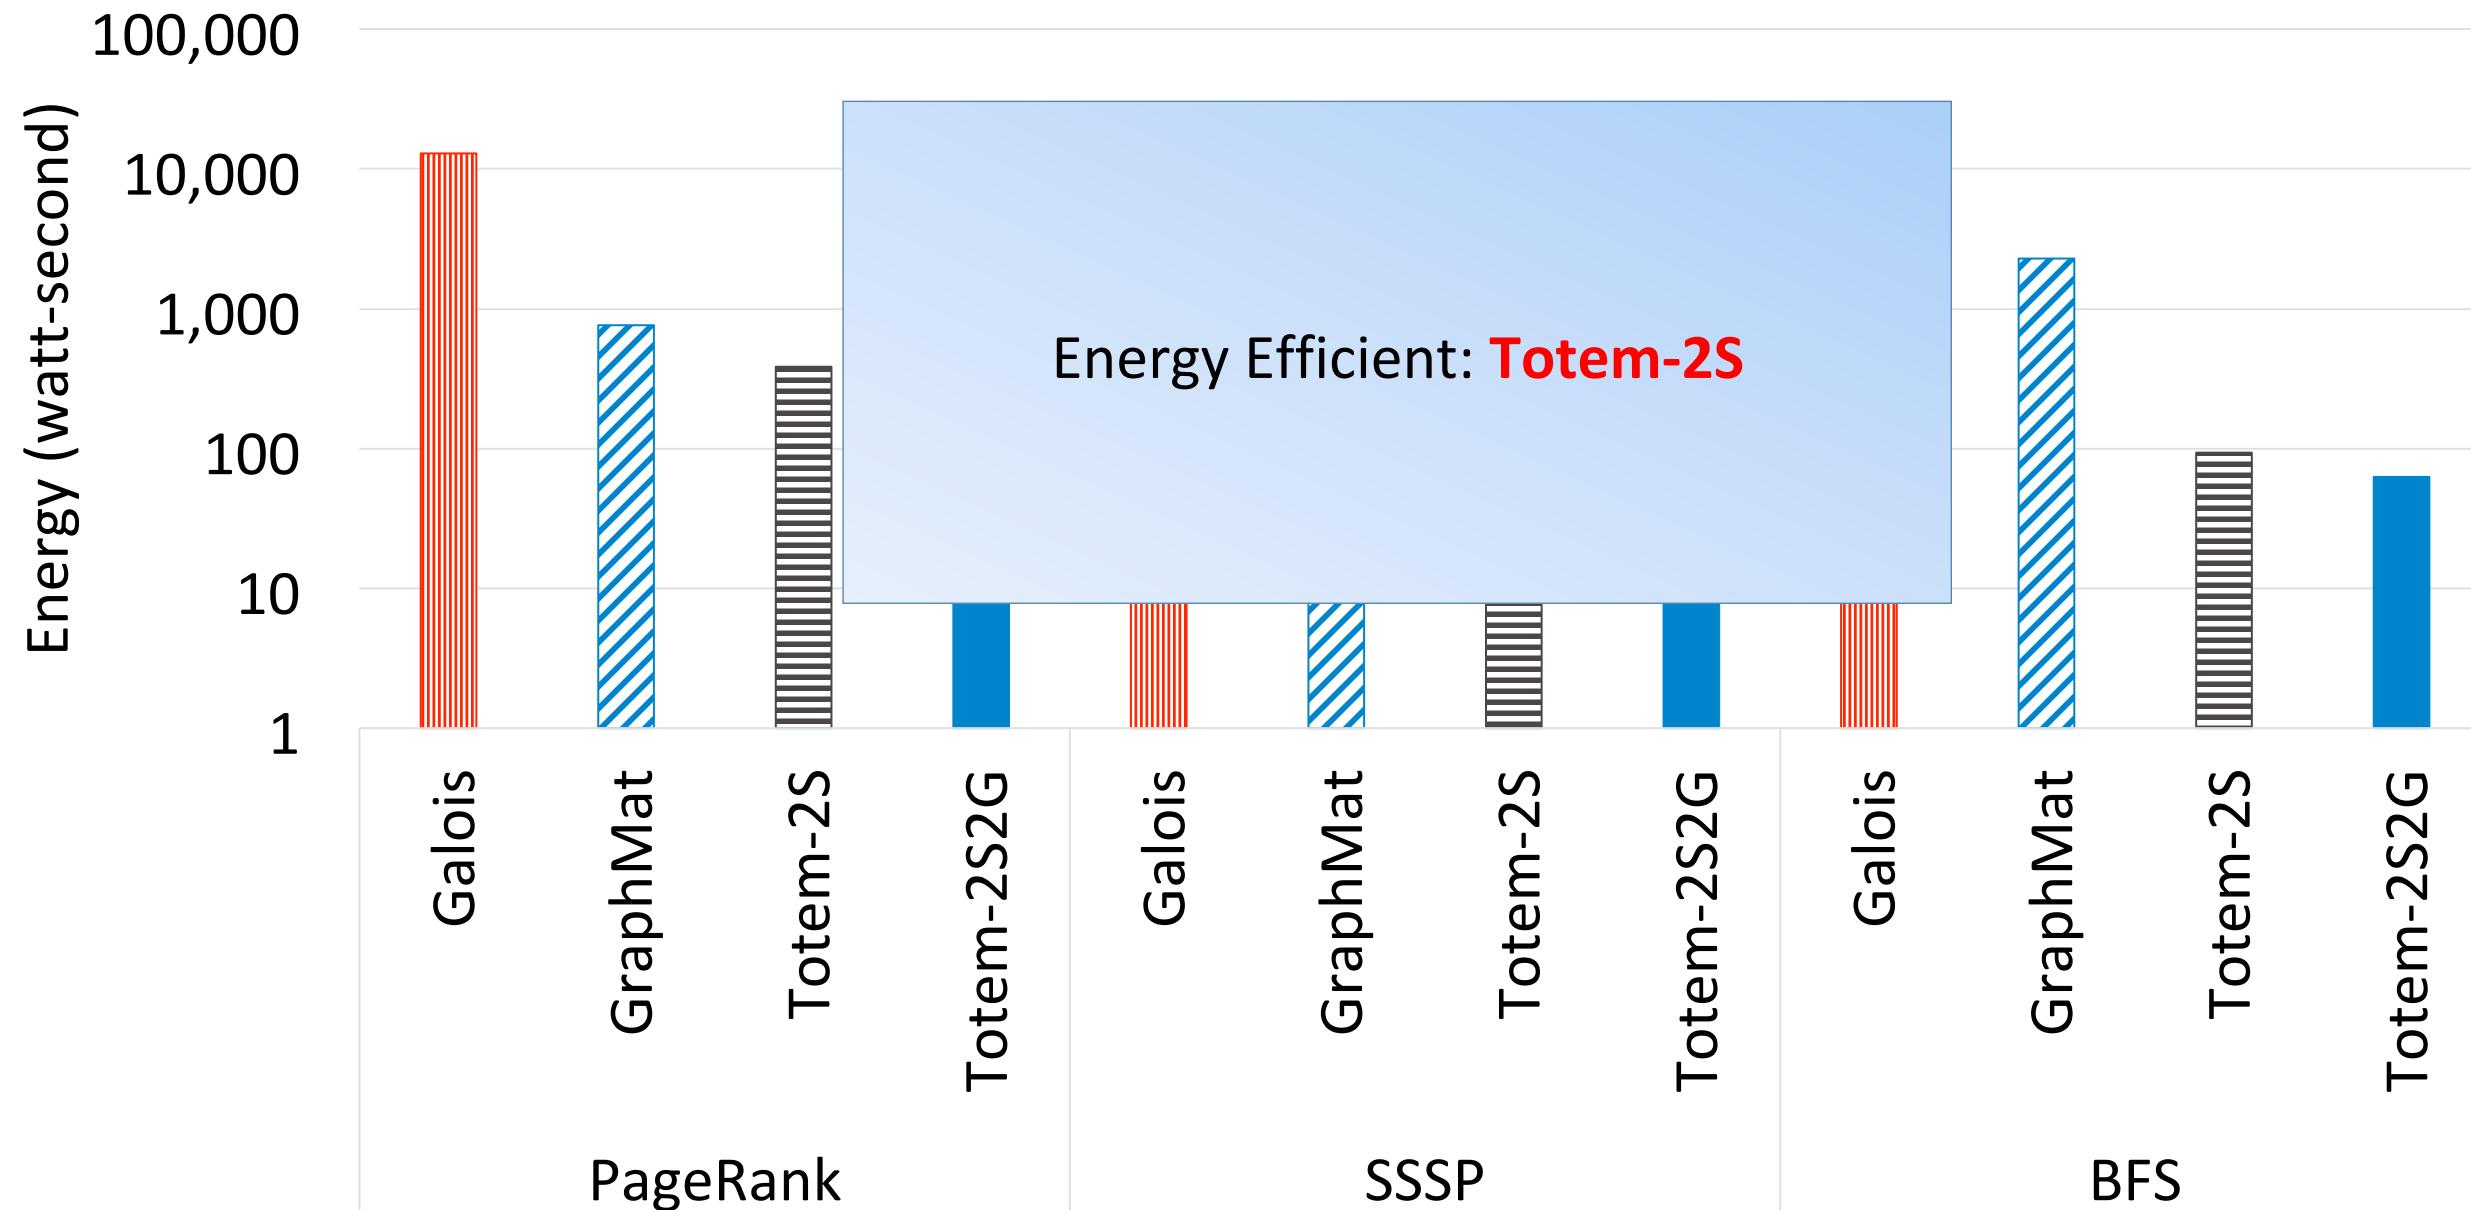
# Experimental Results

## 2. Energy Consumption – CPU Frameworks – Twitter Workload



# Experimental Results

## 2. Energy Consumption – CPU Frameworks – Twitter Workload

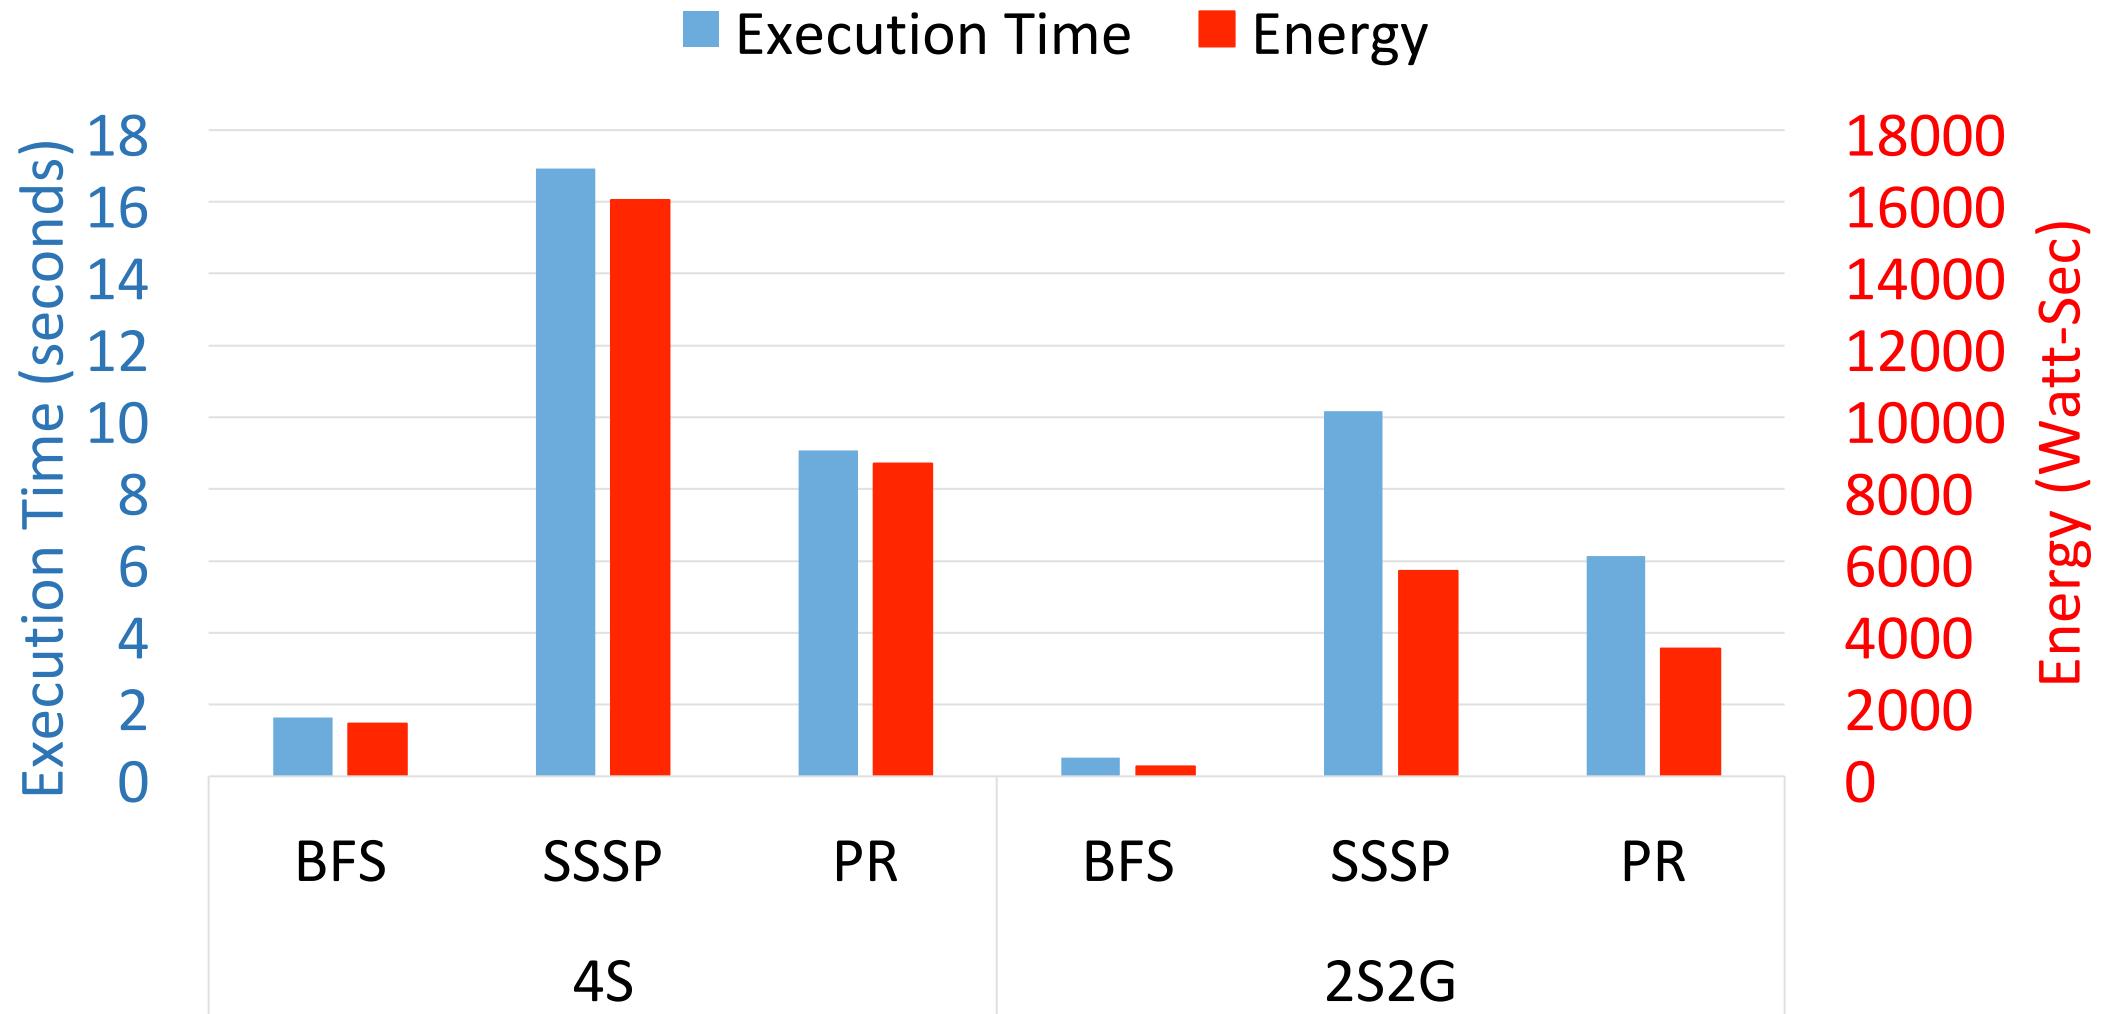


# Summary

- GPU + Linear Algebra | CPU + Vertex programming = Good Match
- GPU based frameworks: ?
- CPU based frameworks: Totem-2S
- Totem Hybrid: Greenest
- CSC  $\longleftrightarrow$  PageRank
- CSR  $\longleftrightarrow$  BFS, SSSP

# Discussion

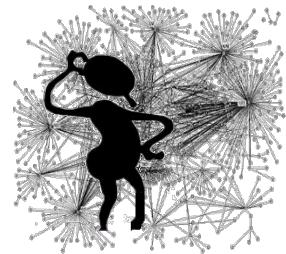
# Does hybrid have the future potential?



**Totem-4S vs Totem-2S2G for RMAT30 (edge list size: 128 GB)**

4S Machine: 4x Intel Xeon E7-4870 v2 (Ivy bridge), with 1,536 GB memory

# Hybrid Graph Processing



## Graph Processing

*Poor locality*

*Data-dependent memory access patterns*

*Low compute-to-memory access ratio*

*Large memory footprint*

*Varying degrees of parallelism  
(both intra- and inter- stage)*

## CPUs



Large Caches +  
summary data  
structures



High Degree

## GPUs



Caches + summary data  
structures

Massive hardware  
multithreading



Low Degree

# Questions

code@:[netsyslab.ece.ubc.ca](mailto:netsyslab.ece.ubc.ca)