

# A Partition-centric Distributed Algorithm for Identifying Euler Circuits in Large Graphs

**Siddharth D Jaiswal** & Yogesh Simmhan

Department of Computational and Data Science,  
Indian Institute of Science, Bangalore, India



2019 IEEE International Workshop on High-Performance Big Data, Deep Learning, and Cloud Computing



# Introduction and Algorithm

---

- Euler Circuits and Use Cases
- Preliminaries
- Algorithmic walkthrough



# Introduction


## Eulerian Circuits-

A circuit starting at some vertex  $v_i$  in a graph such that **every edge** is visited **exactly once** and the circuit ends at the same vertex,  $v_i$ .

These circuits find uses in Biological Networks, Transportation Networks, IoT applications and electrical circuit design.



# Example Eulerian Circuit



All vertices are even degree

A potential Eulerian Circuit-

$(v_1, v_2, v_3, v_4, v_5, v_3, v_{13}, v_{12}, v_{10}, v_9, v_8, v_7, v_6, v_{11}, v_{12}, v_{14}, v_1)$



# Introduction

Existing Parallel Algorithms are **impractical** for commodity clusters and VMs.

We propose an algorithm which **identifies partial paths** and **cycles** within partitions of a large graph, iteratively over multiple Supersteps, using a **Bulk Synchronous Parallel** model.



# Partition Centric Euler Circuit

To find an Eulerian Circuit on a large distributed graph using a **partition-centric** model of computation.

- Undirected Eulerian graph partitioned into **P partitions**.
- A partition has **odd(OB)** and **even(EB)** degree Boundary vertices and **even** degree Internal vertices(**IV**).
- Meta graph-
  - Each **partition** is a meta vertex.
  - All **edges between a pair of partitions** compose a single meta edge.




# Partition Centric Euler Circuit

- Identify pairwise edge-disjoint partial paths locally in each partition.
- Recursively merge the partial paths in pairs of partitions to coarsen them.
- Construct the final circuit



# Partition Centric Euler Circuit



*Partitions:  $P_1$ - $P_4$*

*OB:  $v_1, v_2, v_6, v_9, v_{10}, v_{11}, v_{13}, v_{14}$*

*EB:  $v_3$*

*IV:  $v_4, v_5, v_7, v_8$*


*Local Edges(in Red)*

*Remote Edges(in Blue)*





# Partition Centric Euler Circuit-Meta Graph View





*Meta Vertices:*  $\{MV_1, MV_2, MV_3, MV_4\}$

*Meta Edges:*  $\{ME_1, ME_2, ME_3, ME_4\}$





# Partition Centric Euler Circuit







# Partition Centric Euler Circuit







# Partition Centric Euler Circuit







# Partition Centric Euler Circuit





# Partition Centric Euler Circuit





# Complexity Measures

---

- Coordination
- Computation
- Communication



# Coordination Cost

Number of Supersteps-

Path Identification within a partition - 1  
Superstep

Levels of the Binary Merge Tree-  $\log(n)$

$$\log(n) + 1$$

for  $n$  partitions.





# Computation Cost

Within a partition at every level-

Time to compute the partial paths(cycles) from OB, EB and (if required) IV vertices.

$$O(|BV_i| + |IV_i| + |LE_i|)$$

where,  $BV_i = OB_i + EB_i$  and  $LE_i = \text{Local Edges}$ .



## Communication Cost

Merging of pairs of partitions at every level-  
OB + EB + OB Paths + Remote Edges compose  
this message exchange between partitions.

$$O(|BV_i| + |RE_i|)$$

where,  $BV_i = OB_i + EB_i$  and  $RE_i = \text{Remote Edges}$



# Results

---

- Experimental Setup
- Graph Characteristics
- Discussion

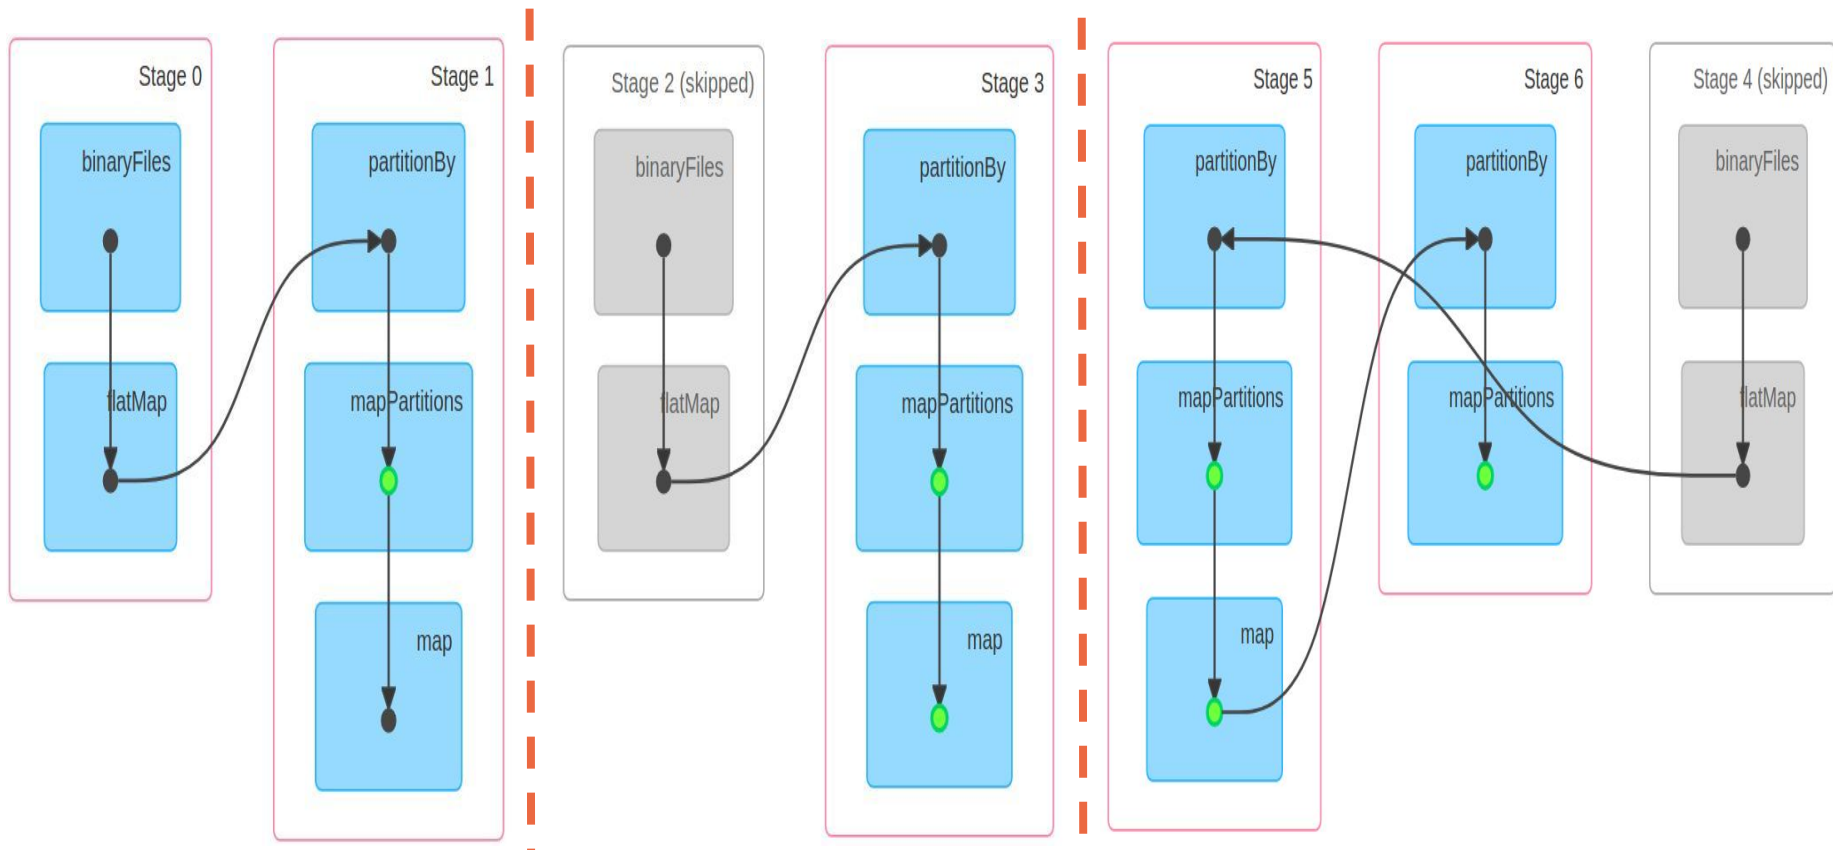


## Results- Experimental Setup

- Apache Spark 2.2 with Java 8
- 8 Microsoft Azure E8/v3 VMs.
  - Each has 8 cores(Intel Xeon E5-2673, 2.3 GHz)
  - 64 GB RAM, 128 GB Local Storage
- 1 Spark Executor/VM assigned 45 GB RAM
- 1 Spark Executor/Partition of the Graph
- Graph Generator Tool: **PaRMat**
- Partitioning Tool: **ParHip**
- Extra Edges: **~5%**



# Apache Spark Pipeline






# Results- Graph Characteristics

<b>Graph</b>	$ V $	$ E $	$\sum_{i \in n}  B_i $	<b>Parts</b> ( $n$ )	$\sum \frac{ R_i }{ E } \%$	$ V_i $	<b>Imbal.%</b>
G20/P2	20M	212M	13M	2	38%		19%
G30/P3	30M	318M	22M	3	49%		48%
G40/P4	40M	423M	23M	4	59%		46%
G40/P8	40M	423M	33M	8	70%		41%
G50/P8	49M	529M	40M	8	70%		63%



# Results- Total Time




- Weak Scaling is **inefficient!**
- Strong Scaling is **inefficient!**
- Tradeoff between compute and, coordination (more supersteps) & data transfer.
- **User compute time** is **~50%** of total time.

**Spark's overheads dominate.**



# Results- Split of User Compute Time- G50/P8



**Level 0-** Object Creation time dominates (~**67%** of total time).

**Level 1-** (De)serialising + Object creation + Merging of partitions dominates.

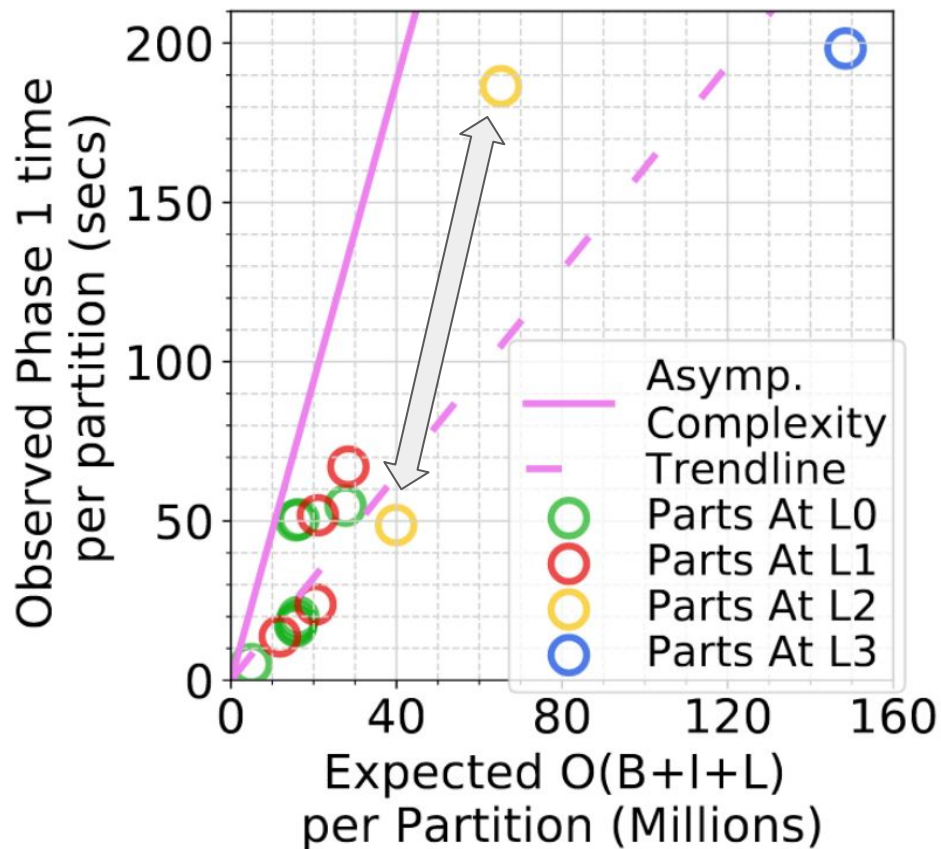
**Level 2 and 3-** Phase 1 dominates (~**48%** and ~**51%** resp).

**Overall compute time increases with level.**






# Results- Expected and Observed Phase 1 Time- G50/P8



- Expected time matches observed time with few outliers.
- Outliers have skew in data.



# Results- Cumulative and Average Memory- G50/P8

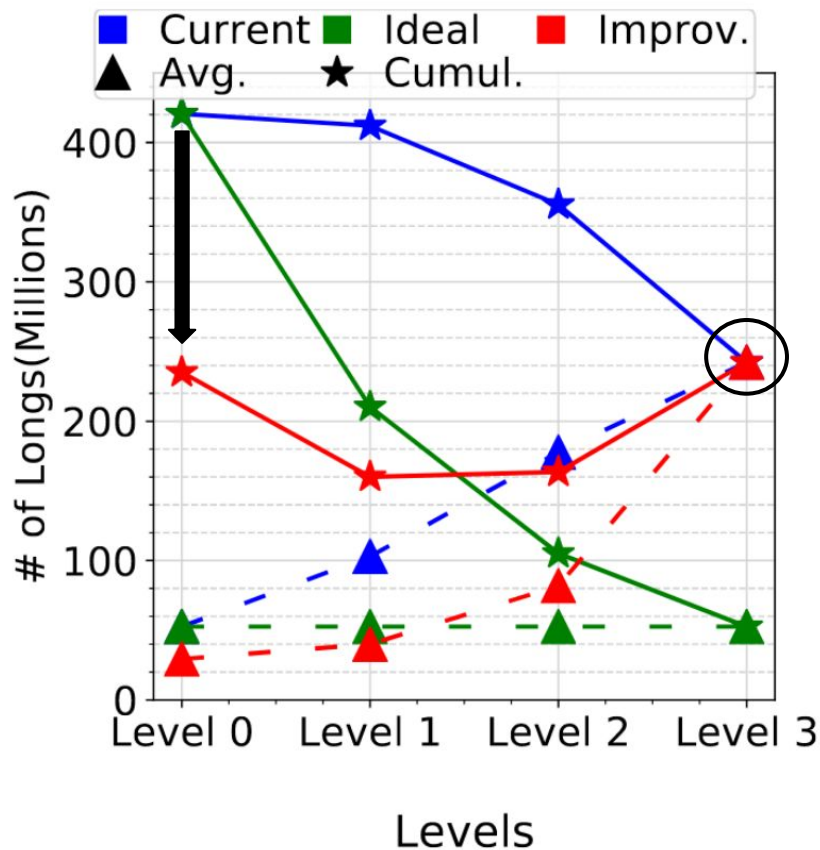


Monotonic drop in cumulative state. Rate of drop is **INADEQUATE**.

Cumulative Drop % is higher with increasing level.



# Results- Cumulative and Average Memory- G50/P8



Two techniques which can **improve memory bottleneck**-

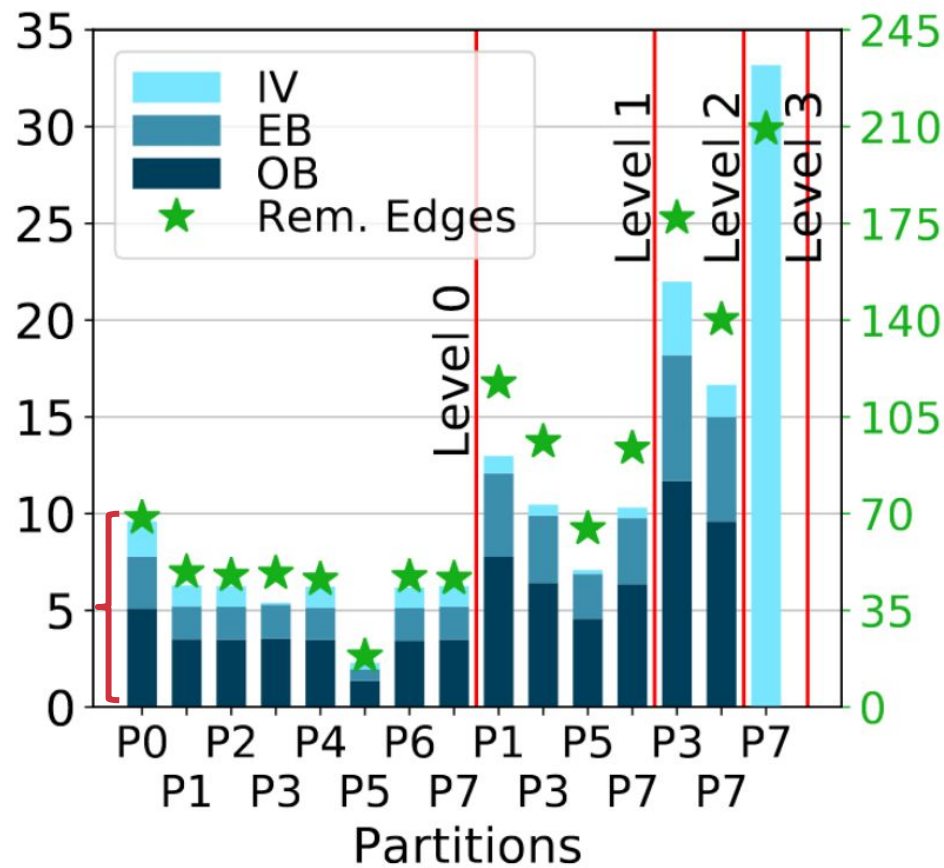
- **Avoiding Duplication** of Remote Edges
- **Defer Transfer** of Remote Edges

At **Level 0**, Initial Memory Maintained drops by **~43%**.

At **Level 3**, Memory remains same as current implementation because all Edges are Local.



# Results- Vertices and Edges per Partition- G50/P8



- Number of boundary vertices and Remote edges grow with merging.
- Remote edge count is **~7x** vertex count.
- Drop in state(IV and LE) is low. IV is a small fraction of total vertices.



# References

1. Jaiswal, Siddharth D., and Yogesh Simmhan. "A Partition-centric Distributed Algorithm for Identifying Euler Circuits in Large Graphs." *arXiv preprint arXiv:1903.06950* (2019).
2. Euler, Leonhard. "Solutio problematis ad geometriam situs pertinentis." *Commentarii academiae scientiarum Petropolitanae* (1741): 128-140.
3. Atallah, Mikhail, and Uzi Vishkin. "Finding Euler tours in parallel." *Journal of Computer and System Sciences* 29.3 (1984): 330-337.
4. Makki, S. A. M. "A distributed algorithm for constructing an Eulerian tour." 1997 IEEE International Performance, Computing and Communications Conference. IEEE, 1997.
5. Awerbuch, Baruch, Amos Israeli, and Yossi Shiloach. "Finding Euler circuits in logarithmic parallel time." *Proceedings of the sixteenth annual ACM symposium on Theory of computing*. ACM, 1984.

# Thank You! Questions?



Contact- [siddharthj@iisc.ac.in](mailto:siddharthj@iisc.ac.in)

Checkout- <http://dream-lab.cds.iisc.ac.in/>  
<https://github.com/dream-lab/>

